

1. Data preprocessing

In [1]:

```
import imageio
import random
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import glob
import cv2
import tensorflow as tf
import tensorflow.keras.layers as tfl

from tensorflow.keras import layers, models
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.layers import Activation, Input, Conv2D, MaxPooling2D, BatchNormalization, Conv2DTranspose, concatenate, Concatenate, UpSampling2D, AveragePooling2D
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.applications import ResNet50
from sklearn.model_selection import train_test_split
import tensorflow.keras as keras

import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
images_list = []
masks_list = []
def store_images(data):
    image = glob.glob('/kaggle/input/lyft-udacity-challenge/' + data + '/' +
data + '/CameraRGB/*.png')
    mask = glob.glob('/kaggle/input/lyft-udacity-challenge/' + data + '/' + d
ata + '/CameraSeg/*.png')
    images_list.extend(image)
    masks_list.extend(mask)
for data in ['dataA', 'dataB', 'dataC', 'dataD', 'dataE']:
    store_images(data)

dataset = tf.data.Dataset.from_tensor_slices((images_list, masks_list))
dataset_length = sum(1 for _ in dataset)

print("Length of dataset:", dataset_length)
# for image, mask in dataset.take(1):
#     print("Image File:", image.numpy().decode(), "Mask File:", mask.nu
mpy().decode())
```

Length of dataset: 5000

In [3]:

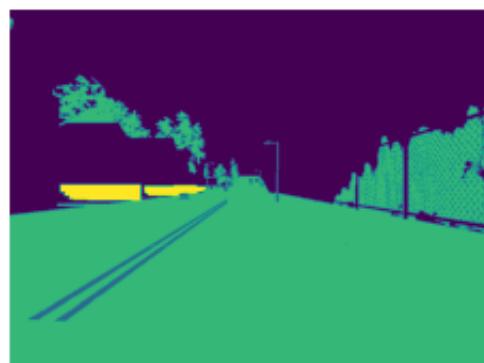
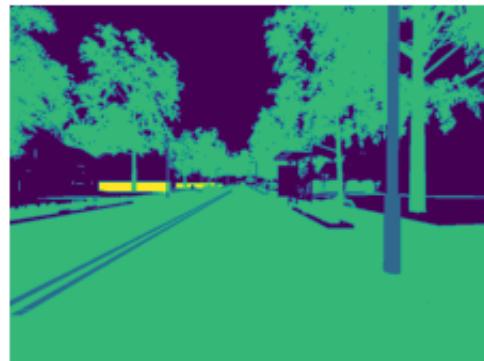
```
def show_images(data, num):
    # Load images
    images = glob.glob('/kaggle/input/lyft-udacity-challenge/' + data + '/' + data + '/CameraRGB/*.png')[ :num]
    masks = glob.glob('/kaggle/input/lyft-udacity-challenge/' + data + '/' + data + '/CameraSeg/*.png')[ :num]

    for i in range(len(images)):
        images[i] = cv2.imread(images[i])
        masks[i] = cv2.imread(masks[i], cv2.IMREAD_GRAYSCALE)

    fig, axes = plt.subplots(num, 2)

    for i in range(len(images)):
        axes[i][0].imshow(cv2.cvtColor(images[i], cv2.COLOR_BGR2RGB))
        axes[i][1].imshow(masks[i])
        axes[i][0].axis('off')
        axes[i][1].axis('off')
    plt.show()

show_images('dataB', 2)
```



In [4]:

```
def load_datas(image_path, mask_path):
    image = tf.io.read_file(image_path)
    image = tf.image.decode_png(image, channels=3)
    image = tf.image.convert_image_dtype(image, tf.float32)

    mask = tf.io.read_file(mask_path)
    mask = tf.image.decode_png(mask, channels=3)
    mask = tf.math.reduce_max(mask, axis=-1, keepdims=True)
    return image, mask

def preprocess(image, mask, height=256, width=256):
    input_image = tf.image.resize(image, (height, width), method='nearest')
    input_mask = tf.image.resize(mask, (height, width), method='nearest')

    return input_image, input_mask

dataset_load = dataset.map(load_datas)
dataset_processed = dataset_load.map(preprocess)
```

In [5]:

```
EPOCHS = 15
BATCH_SIZE = 32
```

In [6]:

```
def split_dataset(data, train_split=0.7, val_split=0.15, test_split=0.15, shuffle=True, shuffle_size=1000, batch_size=BATCH_SIZE):
    """
        Splits the dataset into training, validation, and test sets.
        Returns: (train_dataset, val_dataset, test_dataset).
    """

    assert train_split + val_split + test_split == 1, "Split ratios must sum to 1"
```

```
# Determine the size of the dataset
dataset_size = len(list(data))

# Calculate split sizes
train_size = int(train_split * dataset_size)
val_size = int(val_split * dataset_size)

# Shuffle the dataset
if shuffle:
    data = data.shuffle(shuffle_size, reshuffle_each_iteration=False)

# Split the dataset
train_dataset = data.take(train_size)
val_dataset = data.skip(train_size).take(val_size)
test_dataset = data.skip(train_size + val_size)

# Batch the datasets
train_dataset = train_dataset.batch(batch_size)
val_dataset = val_dataset.batch(batch_size)
test_dataset = test_dataset.batch(batch_size)

return train_dataset, val_dataset, test_dataset

train_dataset, val_dataset, test_dataset = split_dataset(dataset_processed)
```

2. Model Architectures

In this project, 3 model architectures are built, trained and compared, to see each performance on semantic segmentation tasks for autonomous driving.

1. Fully Convolutional Network (FCN)

FCN is a deep learning architecture for semantic segmentation tasks, retaining spatial information with convolutional layers instead of fully connected layers.

2. U-Net

U-Net is a CNN architecture initially used for biomedical image segmentation, featuring a contracting and expansive path with skip connections for precise localization.

3. DeepLabV3

DeepLabV3, developed by Google Research, is a cutting-edge architecture for semantic image segmentation tasks.

- ### Model architecture for FCN

In [7]:

```
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization, ReLU, Conv2DTranspose, concatenate, Activation

def conv_block_3(inputs=None, n_filters=32, dropout_prob=0):
    layer = Conv2D(n_filters, 3, padding='same', kernel_initializer='he_normal')(inputs)
    layer = BatchNormalization(axis=3)(layer)
    layer = Activation("relu")(layer)
    layer = Conv2D(n_filters, 3, padding='same', kernel_initializer='he_normal')(layer)
    layer = BatchNormalization(axis=3)(layer)
    layer = Activation("relu")(layer)
```

```
if dropout_prob > 0:
    layer = tf.keras.layers.Dropout(dropout_prob)(layer)

return layer

def decoder_block(input_tensor, concat_tensor, n_filters):
    # Upsample the input
    decoder = Conv2DTranspose(n_filters, (3, 3), strides=(1, 1), padding='same')(input_tensor)
    # Concatenate with the corresponding layer from the encoder
    decoder = concatenate([decoder, concat_tensor], axis=-1)
    # Additional convolutional layers
    decoder = Conv2D(n_filters, (3, 3), padding='same')(decoder)
    decoder = BatchNormalization()(decoder)
    decoder = Activation('relu')(decoder)

    return decoder

def fcn_model(input_size=(256, 256, 3), n_filters=32, n_classes=13):
    inputs = Input(input_size)

    # Encoder
    cblock1 = conv_block_3(inputs, n_filters) # (None, 256, 256, 32)
    cblock2 = conv_block_3(cblock1, n_filters * 2) # (None, 256, 256, 64)
    cblock3 = conv_block_3(cblock2, n_filters * 4) # (None, 256, 256, 128)

    # Decoder
    dblock1 = decoder_block(cblock3, cblock2, n_filters * 2) # (None, 256, 256, 64)
    dblock2 = decoder_block(dblock1, cblock1, n_filters) # (None, 256, 256, 32)

    # Output layer: (None, 256, 256, 13)
    layer = Conv2D(n_classes, 3, padding='same', kernel_initializer='he
```

```
_normal')(cblock3)
    layer = BatchNormalization(axis=3)(layer)

    model = tf.keras.Model(inputs=inputs, outputs=layer)
    model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(
                      from_logits=True),
                  metrics=['accuracy'])

return model
```

- ### Model architecture for U-Net

In [8]:

```
def conv_layer(inputs, n_filters, kernel_size=3):
    """
    Apply convolution, batch normalization, and ReLU activation.

    Args:
        - inputs: Input tensor.
        - n_filters: Number of filters for the convolutional layer.
        - kernel_size: Size of the convolution kernel.

    Returns:
        - Output tensor after the operations.
    """

    layer = tf1.Conv2D(n_filters, kernel_size=kernel_size, padding='same',
                      kernel_initializer='he_normal')(inputs)
    layer = tf1.BatchNorm(3)(layer)
    layer = tf1.ReLU()(layer)
    return layer
```

In [9]:

```
def conv_block(inputs, n_filters, dropout_prob=0, max_pooling=True):
    """
    Perform two convolutions with optional dropout and max pooling.

    Args:
        - inputs: Input tensor.
        - n_filters: Number of filters for the convolutional layers.
        - dropout_prob: Dropout rate.
        - max_pooling: Boolean, whether to include a max pooling layer.

    Returns:
        - next_layer: Output tensor for the next layer.
        - skip_connection: Output tensor for the skip connection.
    """

    # Two convolutional layers
    layer = conv_layer(inputs, n_filters)
    layer = conv_layer(layer, n_filters)

    # Optional dropout
    if dropout_prob > 0:
        layer = tf1.Dropout(dropout_prob)(layer)
    # Optional max pooling
    if max_pooling:
        next_layer = tf1.MaxPooling2D(pool_size=(2, 2))(layer)
    else:
        next_layer = layer

    skip_connection = layer

    return next_layer, skip_connection
```

In [10]:

```
def upsampling_block(inputs, skip_connection_inputs, n_filters):  
    """  
        Upsample the input and merge with the skip connection.  
  
    Args:  
        - inputs: Input tensor from the previous layer.  
        - skip_connection_inputs: Input tensor from the corresponding contraction block (for the skip connection).  
        - n_filters: Number of filters for the convolutional layers.  
  
    Returns:  
        - layer: Output tensor after upsampling and convolution.  
    """  
  
    # Upsampling  
    up = tf1.Conv2DTranspose(n_filters, kernel_size=3, strides=(2, 2),  
                           padding='same')(inputs)  
  
    # Merging with skip connection  
    merge = tf1.concatenate([up, skip_connection_inputs], axis=3)  
  
    # Two convolutional layers  
    layer = conv_layer(merge, n_filters)  
    layer = conv_layer(layer, n_filters)  
  
    return layer
```

In [11]:

```
# U-Net model  
def unet_model(input_size=(256, 256, 3), n_filters=32, n_classes=13):  
    """  
        Define the U-Net model architecture.  
  
    Args:  
        - input_size: Shape of the input images.  
        - n_filters: Number of filters for the convolutional layers in the first block. Gets doubled in each subsequent block.  
        - n_classes: Number of output classes.
```

```
Returns:
- model: Compiled U-Net model.
"""

inputs = tf1.Input(input_size)

# Encoding path
cblock1 = conv_block(inputs, n_filters)
cblock2 = conv_block(cblock1[0], n_filters * 2)
cblock3 = conv_block(cblock2[0], n_filters * 4)
cblock4 = conv_block(cblock3[0], n_filters * 8, dropout_prob=0.3)
cblock5 = conv_block(cblock4[0], n_filters * 16, dropout_prob=0.3,
max_pooling=False)

# Decoding path
ublock6 = upsampling_block(cblock5[0], cblock4[1], n_filters * 8)
ublock7 = upsampling_block(ublock6, cblock3[1], n_filters * 4)
ublock8 = upsampling_block(ublock7, cblock2[1], n_filters * 2)
ublock9 = upsampling_block(ublock8, cblock1[1], n_filters)

# Output layer
output_layer = conv_layer(ublock9, n_filters)
output_layer = tf1.Conv2D(n_classes, kernel_size=1, padding='same')
(output_layer)

model = tf.keras.Model(inputs=inputs, outputs=output_layer)
model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

return model
```

- ### Model architecture for DeepLabV3

In [12]:

```
def aspp(input):
    shape = input.shape
    y_pool = AveragePooling2D(pool_size = (shape[1], shape[2]))(input)
    y_pool = Conv2D(filters=256, padding='same', use_bias=False, kernel_size=1)(y_pool)
    y_pool = BatchNormalization()(y_pool)
    y_pool = Activation("relu")(y_pool)
    y_pool = UpSampling2D((shape[1], shape[2]), interpolation="bilinear")(y_pool)

    out_1 = Conv2D(filters=256, padding='same', dilation_rate = 1, use_bias=False, kernel_size=1)(input)
    out_1 = BatchNormalization()(out_1)
    out_1 = Activation("relu")(out_1)

    out_6 = Conv2D(filters=256, padding='same', dilation_rate = 6, use_bias=False, kernel_size=1)(input)
    out_6 = BatchNormalization()(out_6)
    out_6 = Activation("relu")(out_6)

    out_12 = Conv2D(filters=256, padding='same', dilation_rate = 12, use_bias=False, kernel_size=1)(input)
    out_12 = BatchNormalization()(out_12)
    out_12 = Activation("relu")(out_12)

    out_18 = Conv2D(filters=256, padding='same', dilation_rate = 18, use_bias=False, kernel_size=1)(input)
    out_18 = BatchNormalization()(out_18)
    out_18 = Activation("relu")(out_18)

    y = Concatenate()([y_pool, out_1, out_6, out_12, out_18])
    y = Conv2D(filters=256, padding='same', dilation_rate = 1, use_bias=False, kernel_size=1)(y)
    y = BatchNormalization()(y)
    y = Activation("relu")(y)
```

```
    return y

def DeeplabV3(image_size=(256, 256, 3), num_classes=13):

    inputs = Input(image_size)

    resnet50 = keras.applications.ResNet50(
        weights="imagenet", include_top=False, input_tensor=inputs
    )

    # pretrained Resnet50 output
    # x1 = resnet50.get_layer("conv4_block6_out").output
    x1 = resnet50.get_layer("conv4_block6_2_relu").output
    x1 = aspp(x1)
    x1 = UpSampling2D((4, 4), interpolation="bilinear")(x1)

    # low level features
    # x2 = resnet50.get_layer("conv2_block2_out").output
    x2 = resnet50.get_layer("conv2_block3_2_relu").output
    x2 = Conv2D(filters=48, padding='same', use_bias=False, kernel_size=1)(x2)
    x2 = BatchNormalization()(x2)
    x2 = Activation("relu")(x2)

    x = Concatenate()([x1, x2])

    x = Conv2D(filters=256, padding='same', activation='relu', use_bias=False, kernel_size=3)(x)
    x = BatchNormalization()(x)
    x = Activation("relu")(x)

    x = Conv2D(filters=256, padding='same', activation='relu', use_bias=False, kernel_size=3)(x)
    x = BatchNormalization()(x)
    x = Activation("relu")(x)

    x = UpSampling2D((4, 4), interpolation="bilinear")(x)
```

```
# output
x = Conv2D(num_classes, kernel_size=(1, 1), name="output_layer", padding="same")(x)
x = Activation('sigmoid')(x)

model = Model(inputs=inputs, outputs=x)

model.compile(optimizer=tf.keras.optimizers.Adam(),
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

return model
```

3. Model Training

Hyperparameter tuning and model configurations.

In [13]:

```
img_height = 256
img_width = 256
num_channels = 3
filters = 32
n_classes = 13
```

Function to plot training history

In [14]:

```
def plot_training_history(history):
    acc = [0.] + history.history.get('accuracy', [])
    val_acc = [0.] + history.history.get('val_accuracy', [])
    loss = history.history.get('loss', [])
    val_loss = history.history.get('val_loss', [])

    # Plotting
    plt.figure(figsize=(8, 8))

    # Accuracy subplot
    plt.subplot(2, 1, 1)
    plt.plot(acc, label='Training Accuracy')
    plt.plot(val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.ylabel('Accuracy')
    plt.title('Training and Validation Accuracy')
    plt.ylim([0, 1])

    # Loss subplot
    plt.subplot(2, 1, 2)
    plt.plot(loss, label='Training Loss')
    plt.plot(val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.ylabel('Cross Entropy')
    plt.title('Training and Validation Loss')
    plt.xlabel('Epoch')
    plt.ylim([0, 1])

    plt.tight_layout()
    plt.show()
```

- ### Training FCN

In [15]:

```
import time

# FCN model

fcn_model1 = fcn_model()
fcn_model1.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param
input_layer (InputLayer)	(None, 256, 256, 3)	
conv2d (Conv2D)	(None, 256, 256, 32)	896
batch_normalization (BatchNormalization)	(None, 256, 256, 32)	128
activation (Activation)	(None, 256, 256, 32)	
conv2d_1 (Conv2D)	(None, 256, 256, 32)	9,264
batch_normalization_1 (BatchNormalization)	(None, 256, 256, 32)	128
activation_1 (Activation)	(None, 256, 256, 32)	
conv2d_2 (Conv2D)	(None, 256, 256, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 256, 256, 64)	256
activation_2 (Activation)	(None, 256, 256, 64)	
conv2d_3 (Conv2D)	(None, 256, 256, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 256, 256, 64)	256
activation_3 (Activation)	(None, 256, 256, 64)	
conv2d_4 (Conv2D)	(None, 256, 256, 128)	73,856

batch_normalization_4 (BatchNormalization)	(None, 256, 256, 128)	51
activation_4 (Activation)	(None, 256, 256, 128)	
conv2d_5 (Conv2D)	(None, 256, 256, 128)	147,58
batch_normalization_5 (BatchNormalization)	(None, 256, 256, 128)	51
activation_5 (Activation)	(None, 256, 256, 128)	
conv2d_8 (Conv2D)	(None, 256, 256, 13)	14,98
batch_normalization_8 (BatchNormalization)	(None, 256, 256, 13)	5

Total params: 303,841 (1.16 MB)

Trainable params: 302,919 (1.16 MB)

Non-trainable params: 922 (3.60 KB)

In [16]:

```
start_time = time.time()

# reduce_lr = ReduceLROnPlateau(monitor= "loss", factor=0.1,
#                               patience= 1, min_lr= 1e-6)
# early_stop = EarlyStopping(patience= 2)

reduce_lr = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,
                              patience=10, min_lr=1e-6)
early_stop = EarlyStopping(monitor='val_accuracy', patience=5, restore_
best_weights=True)

fcn_result1 = fcn_model1.fit(train_dataset,
                             validation_data=val_dataset,
                             epochs=EPOCHS,
                             callbacks=[reduce_lr,early_stop],
                             batch_size=BATCH_SIZE)

end_time = time.time()
training_time = end_time - start_time
```

Epoch 1/15

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1713213812.688247 68 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

110/110  **220s** 1s/step - accuracy: 0.6441 - loss: 1.5717 - val_accuracy: 0.3936 - val_loss: 2.2648 - learning_rate: 0.0010

Epoch 2/15

110/110  **94s** 778ms/step - accuracy: 0.8108 - loss: 1.0322 - val_accuracy: 0.3834 - val_loss: 2.4630 - learning_rate: 0.010

Epoch 3/15

110/110  **94s** 780ms/step - accuracy: 0.8382 - loss: 0.8588 - val_accuracy: 0.5838 - val_loss: 1.6576 - learning_rate: 0.0010

Epoch 4/15

110/110  **94s** 777ms/step - accuracy: 0.8538 - loss: 0.7404 - val_accuracy: 0.7593 - val_loss: 1.0123 - learning_rate: 0.0010

Epoch 5/15

110/110  **93s** 778ms/step - accuracy: 0.8637 - loss: 0.6558 - val_accuracy: 0.8455 - val_loss: 0.6868 - learning_rate: 0.0010

Epoch 6/15

110/110  **93s** 778ms/step - accuracy: 0.8738 - loss: 0.5846 - val_accuracy: 0.8668 - val_loss: 0.5881 - learning_rate: 0.0010

Epoch 7/15

110/110  **93s** 778ms/step - accuracy: 0.8802 - loss: 0.5330 - val_accuracy: 0.8791 - val_loss: 0.5168 - learning_rate: 0.0010

Epoch 8/15

110/110  **93s** 779ms/step - accuracy: 0.8858 - loss: 0.4913 - val_accuracy: 0.8831 - val_loss: 0.4882 - learning_rate: 0.0010

Epoch 9/15

110/110  **93s** 777ms/step - accuracy: 0.8907 - loss: 0.4567 - val_accuracy: 0.8751 - val_loss: 0.4884 - learning_rate: 0.0010

Epoch 10/15

110/110  **93s** 778ms/step - accuracy: 0.8946 - loss: 0.4286 - val_accuracy: 0.8782 - val_loss: 0.4606 - learning_rate: 0.0010

Epoch 11/15

110/110  **94s** 786ms/step - accuracy: 0.8982 - loss: 0.4043 - val_accuracy: 0.8862 - val_loss: 0.4266 - learning_rate: 0.0010

Epoch 12/15

110/110 ————— **94s** 782ms/step - accuracy: 0.9011 - loss: 0.3847 - val_accuracy: 0.8853 - val_loss: 0.4205 - learning_rate: 0.010
Epoch 13/15
110/110 ————— **94s** 779ms/step - accuracy: 0.9036 - loss: 0.3678 - val_accuracy: 0.8865 - val_loss: 0.4068 - learning_rate: 0.010
Epoch 14/15
110/110 ————— **93s** 778ms/step - accuracy: 0.9061 - loss: 0.3526 - val_accuracy: 0.8967 - val_loss: 0.3737 - learning_rate: 0.010
Epoch 15/15
110/110 ————— **93s** 772ms/step - accuracy: 0.9084 - loss: 0.3390 - val_accuracy: 0.9058 - val_loss: 0.3440 - learning_rate: 0.010

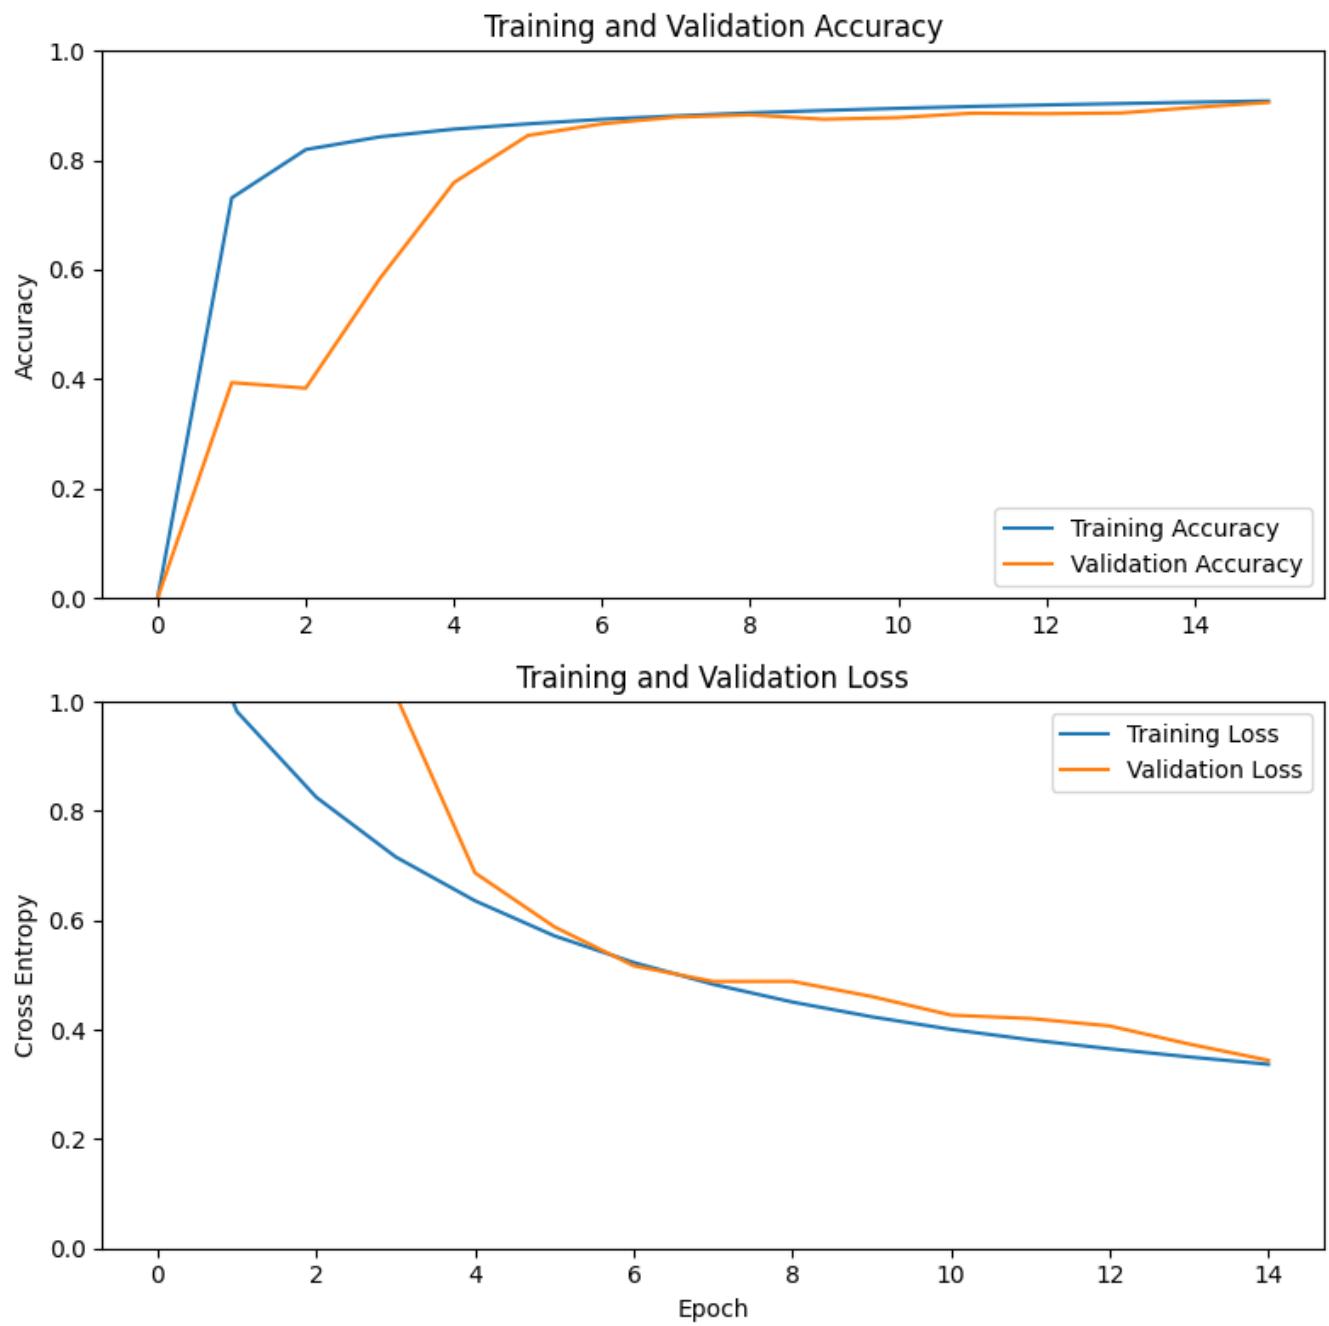
In [17]:

```
print("-----")
print("Training time for FCN:", training_time, "seconds")
```

Training time for FCN: 1530.48957157135 seconds

In [18]:

```
plot_training_history(fcn_result1)
```



- ### Training U-Net

In [19]:

```
# U-net model
unet_model1 = unet_model()
unet_model1.summary()
```

Model: "functional_3"

Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, 256, 256, 3)	0	-
conv2d_9 (Conv2D)	(None, 256, 256, 32)	896	input_layer_1[0]
batch_normalizatio... (BatchNormalizatio...)	(None, 256, 256, 32)	128	conv2d_9[0][0]
re_lu (ReLU)	(None, 256, 256, 32)	0	batch_normalizat
conv2d_10 (Conv2D)	(None, 256, 256, 32)	9,248	re_lu[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 256, 256, 32)	128	conv2d_10[0][0]
re_lu_1 (ReLU)	(None, 256, 256, 32)	0	batch_normalizat
max_pooling2d (MaxPooling2D)	(None, 128, 128, 32)	0	re_lu_1[0][0]
conv2d_11 (Conv2D)	(None, 128, 128, 64)	18,496	max_pooling2d[0]
batch_normalizatio... (BatchNormalizatio...)	(None, 128, 128, 64)	256	conv2d_11[0][0]
re_lu_2 (ReLU)	(None, 128, 128, 64)	0	batch_normalizat
conv2d_12 (Conv2D)	(None, 128, 128, 64)	36,928	re_lu_2[0][0]

batch_normalizatio... (BatchNormalizatio...	(None, 128, 128, 64)	256	conv2d_12[0][0]
re_lu_3 (ReLU)	(None, 128, 128, 64)	0	batch_normalizati...
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 64)	0	re_lu_3[0][0]
conv2d_13 (Conv2D)	(None, 64, 64, 128)	73,856	max_pooling2d_1[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 64, 64, 128)	512	conv2d_13[0][0]
re_lu_4 (ReLU)	(None, 64, 64, 128)	0	batch_normalizati...
conv2d_14 (Conv2D)	(None, 64, 64, 128)	147,584	re_lu_4[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 64, 64, 128)	512	conv2d_14[0][0]
re_lu_5 (ReLU)	(None, 64, 64, 128)	0	batch_normalizati...
max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 128)	0	re_lu_5[0][0]
conv2d_15 (Conv2D)	(None, 32, 32, 256)	295,168	max_pooling2d_2[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 32, 32, 256)	1,024	conv2d_15[0][0]
re_lu_6 (ReLU)	(None, 32, 32, 256)	0	batch_normalizati...
conv2d_16 (Conv2D)	(None, 32, 32, 256)	590,080	re_lu_6[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 32, 32, 256)	1,024	conv2d_16[0][0]
re_lu_7 (ReLU)	(None, 32, 32, 256)	0	batch_normalizati...
dropout (Dropout)	(None, 32, 32, 256)	0	re_lu_7[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 256)	0	dropout[0][0]

conv2d_17 (Conv2D)	(None, 16, 16, 512)	1,180,160	max_pooling2d_3[0][0]
batch_normalization_8 (BatchNormalization)	(None, 16, 16, 512)	2,048	conv2d_17[0][0]
re_lu_8 (ReLU)	(None, 16, 16, 512)	0	batch_normalization_8[0][0]
conv2d_18 (Conv2D)	(None, 16, 16, 512)	2,359,808	re_lu_8[0][0]
batch_normalization_9 (BatchNormalization)	(None, 16, 16, 512)	2,048	conv2d_18[0][0]
re_lu_9 (ReLU)	(None, 16, 16, 512)	0	batch_normalization_9[0][0]
dropout_1 (Dropout)	(None, 16, 16, 512)	0	re_lu_9[0][0]
conv2d_transpose_2 (Conv2DTranspose)	(None, 32, 32, 256)	1,179,904	dropout_1[0][0]
concatenate_2 (Concatenate)	(None, 32, 32, 512)	0	conv2d_transpose_2[0][0]
conv2d_19 (Conv2D)	(None, 32, 32, 256)	1,179,904	concatenate_2[0][0]
batch_normalization_10 (BatchNormalization)	(None, 32, 32, 256)	1,024	conv2d_19[0][0]
re_lu_10 (ReLU)	(None, 32, 32, 256)	0	batch_normalization_10[0][0]
conv2d_20 (Conv2D)	(None, 32, 32, 256)	590,080	re_lu_10[0][0]
batch_normalization_11 (BatchNormalization)	(None, 32, 32, 256)	1,024	conv2d_20[0][0]
re_lu_11 (ReLU)	(None, 32, 32, 256)	0	batch_normalization_11[0][0]
conv2d_transpose_3 (Conv2DTranspose)	(None, 64, 64, 128)	295,040	re_lu_11[0][0]
concatenate_3 (Concatenate)	(None, 64, 64, 256)	0	conv2d_transpose_3[0][0]
conv2d_21 (Conv2D)	(None, 64, 64, 256)	295,040	concatenate_3[0][0]

	128)		
batch_normalization_... (BatchNormalizatio...)	(None, 64, 64, 128)	512	conv2d_21[0][0]
re_lu_12 (ReLU)	(None, 64, 64, 128)	0	batch_normalizati...
conv2d_22 (Conv2D)	(None, 64, 64, 128)	147,584	re_lu_12[0][0]
batch_normalization_... (BatchNormalizatio...)	(None, 64, 64, 128)	512	conv2d_22[0][0]
re_lu_13 (ReLU)	(None, 64, 64, 128)	0	batch_normalizati...
conv2d_transpose_4 (Conv2DTranspose)	(None, 128, 128, 64)	73,792	re_lu_13[0][0]
concatenate_4 (Concatenate)	(None, 128, 128, 128)	0	conv2d_transpose_4[0][0]
conv2d_23 (Conv2D)	(None, 128, 128, 64)	73,792	concatenate_4[0][0]
batch_normalization_... (BatchNormalizatio...)	(None, 128, 128, 64)	256	conv2d_23[0][0]
re_lu_14 (ReLU)	(None, 128, 128, 64)	0	batch_normalizati...
conv2d_24 (Conv2D)	(None, 128, 128, 64)	36,928	re_lu_14[0][0]
batch_normalization_... (BatchNormalizatio...)	(None, 128, 128, 64)	256	conv2d_24[0][0]
re_lu_15 (ReLU)	(None, 128, 128, 64)	0	batch_normalizati...
conv2d_transpose_5 (Conv2DTranspose)	(None, 256, 256, 32)	18,464	re_lu_15[0][0]
concatenate_5 (Concatenate)	(None, 256, 256, 64)	0	conv2d_transpose_5[0][0]
conv2d_25 (Conv2D)	(None, 256, 256, 32)	18,464	concatenate_5[0][0]
batch_normalization_... (BatchNormalizatio...)	(None, 256, 256, 32)	128	conv2d_25[0][0]

re_lu_16 (ReLU)	(None, 256, 256, 32)	0	batch_normalizat
conv2d_26 (Conv2D)	(None, 256, 256, 32)	9,248	re_lu_16[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 256, 256, 32)	128	conv2d_26[0][0]
re_lu_17 (ReLU)	(None, 256, 256, 32)	0	batch_normalizat
conv2d_27 (Conv2D)	(None, 256, 256, 32)	9,248	re_lu_17[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 256, 256, 32)	128	conv2d_27[0][0]
re_lu_18 (ReLU)	(None, 256, 256, 32)	0	batch_normalizat
conv2d_28 (Conv2D)	(None, 256, 256, 13)	429	re_lu_18[0][0]

Total params: 8,652,045 (33.00 MB)

Trainable params: 8,646,093 (32.98 MB)

Non-trainable params: 5,952 (23.25 KB)

In [20]:

```
start_time = time.time()

unet_model1 = unet_model()
reduce_lr = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1,
                               patience=10, min_lr=1e-6)
early_stop = EarlyStopping(monitor='val_accuracy', patience=5, restore_
best_weights=True)

unet_result1 = unet_model1.fit(train_dataset,
                               validation_data=val_dataset,
                               epochs=EPOCHS,
                               callbacks=[reduce_lr, early_stop],
                               batch_size=BATCH_SIZE)
end_time = time.time()
```

Epoch 1/15

110/110  **180s** 1s/step - accuracy: 0.7155 - loss: 1.1250 - val_accuracy: 0.8096 - val_loss: 0.7900 - learning_rate: 0.0010

Epoch 2/15

110/110  **86s** 701ms/step - accuracy: 0.9188 - loss: 0.3165 - val_accuracy: 0.8040 - val_loss: 0.8075 - learning_rate: 0.0010

Epoch 3/15

110/110  **86s** 707ms/step - accuracy: 0.9408 - loss: 0.2140 - val_accuracy: 0.8134 - val_loss: 0.7329 - learning_rate: 0.0010

Epoch 4/15

110/110  **85s** 702ms/step - accuracy: 0.9514 - loss: 0.1691 - val_accuracy: 0.9216 - val_loss: 0.2815 - learning_rate: 0.0010

Epoch 5/15

110/110  **86s** 702ms/step - accuracy: 0.9589 - loss: 0.1393 - val_accuracy: 0.9363 - val_loss: 0.2229 - learning_rate: 0.0010

Epoch 6/15

110/110  **85s** 704ms/step - accuracy: 0.9643 - loss: 0.1176 - val_accuracy: 0.9401 - val_loss: 0.2012 - learning_rate: 0.010

Epoch 7/15

110/110  **85s** 702ms/step - accuracy: 0.9682 - loss: 0.1021 - val_accuracy: 0.9607 - val_loss: 0.1281 - learning_rate: 0.010

Epoch 8/15

110/110  **85s** 700ms/step - accuracy: 0.9698 - loss: 0.0948 - val_accuracy: 0.9672 - val_loss: 0.1028 - learning_rate: 0.010

Epoch 9/15

110/110  **86s** 708ms/step - accuracy: 0.9719 - loss: 0.0873 - val_accuracy: 0.9697 - val_loss: 0.0944 - learning_rate: 0.010

Epoch 10/15

110/110  **85s** 702ms/step - accuracy: 0.9744 - loss: 0.0781 - val_accuracy: 0.9706 - val_loss: 0.0904 - learning_rate: 0.010

Epoch 11/15

110/110  **85s** 702ms/step - accuracy: 0.9764 - loss: 0.0709 - val_accuracy: 0.9723 - val_loss: 0.0842 - learning_rate: 0.010

Epoch 12/15

110/110  **85s** 705ms/step - accuracy: 0.9778 - loss: 0.0660 - val_accuracy: 0.9745 - val_loss: 0.0774 - learning_rate: 0.010

Epoch 13/15

110/110  **85s** 699ms/step - accuracy: 0.9788 - loss: 0.0623 - val_accuracy: 0.9649 - val_loss: 0.1127 - learning_rate: 0.010

Epoch 14/15

110/110  **143s** 706ms/step - accuracy: 0.9679 - loss: 0.0989 - val_accuracy: 0.9740 - val_loss: 0.0789 - learning_rate: 0.0010

Epoch 15/15

110/110 ————— **86s** 709ms/step - accuracy: 0.9786 - loss: 0.0631 - val_accuracy: 0.9773 - val_loss: 0.0682 - learning_rate: 0.0010

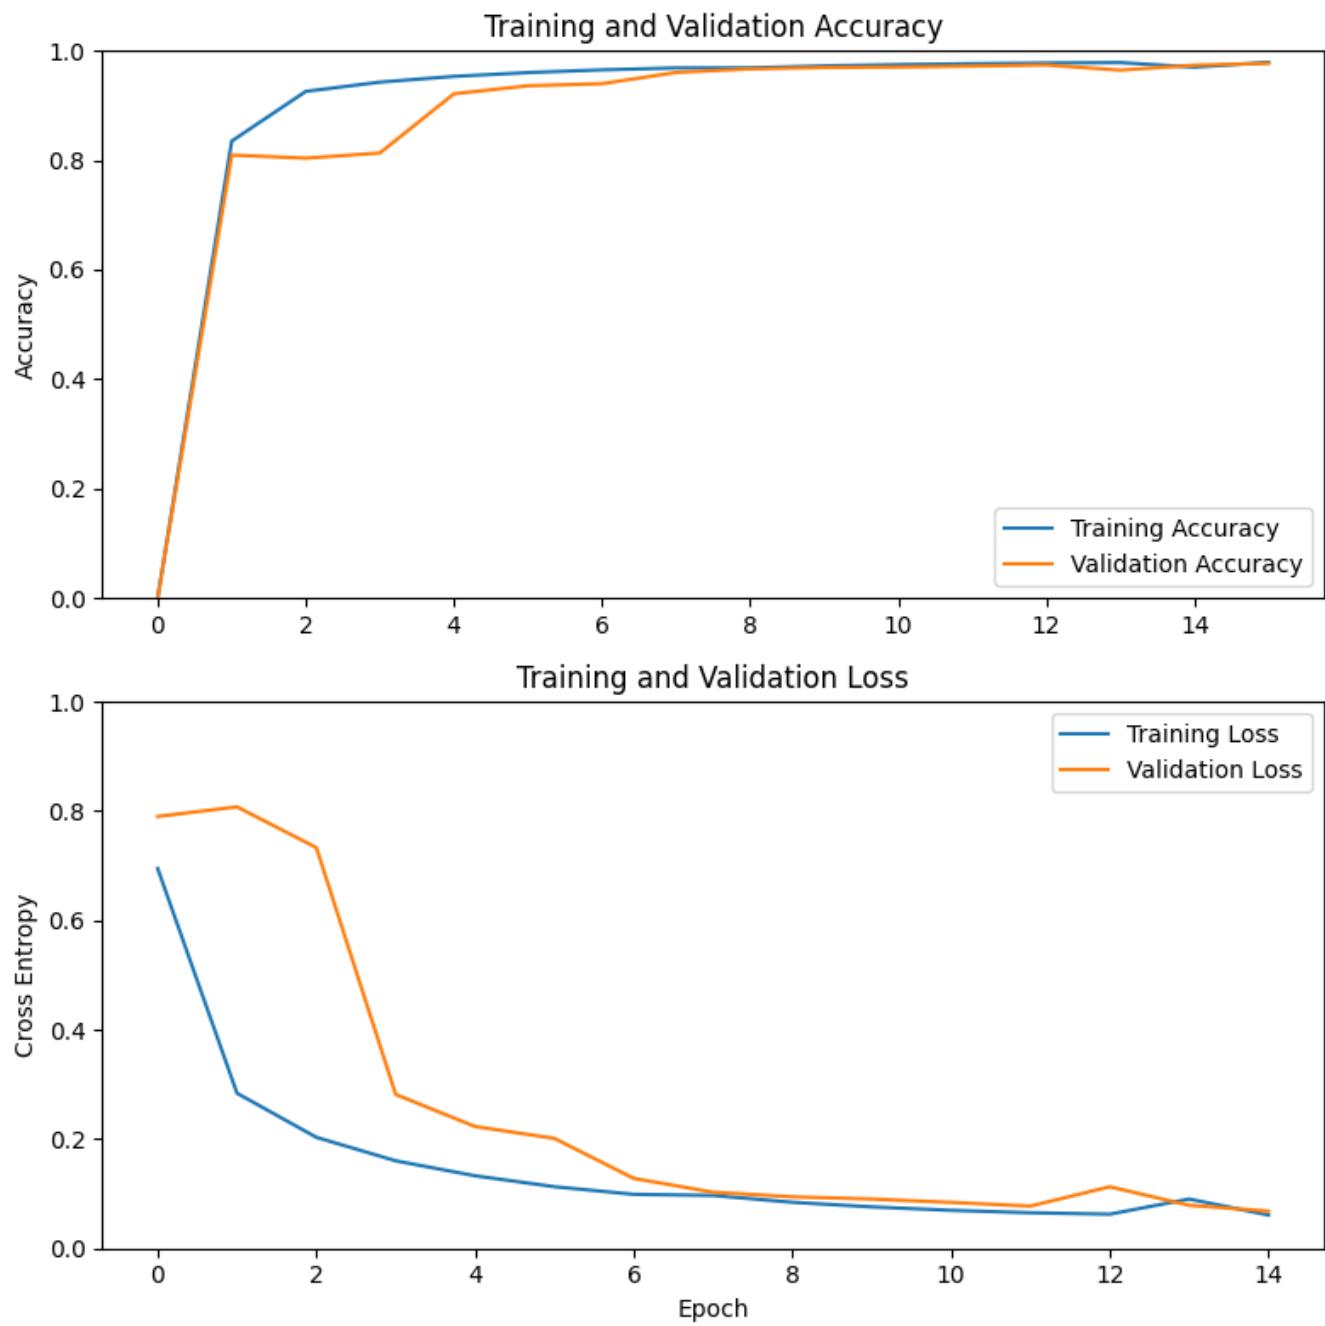
In [21]:

```
training_time = end_time - start_time
print("-----")
print("Training time for U-Net:", training_time, "seconds")
```

Training time for U-Net: 1433.0051529407501 seconds

In [22]:

```
plot_training_history(unet_result1)
```



- ### Training DeepLabV3

In [23]:

```
# Deeplab model
deeplab_model1 = DeeplabV3()
deeplab_model1.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 ————— **1s** 0us/step

Model: "functional_7"

Layer (type)	Output Shape	Param #	Connected to
input_layer_3 (InputLayer)	(None, 256, 256, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 262, 262, 3)	0	input_layer_3[0]
conv1_conv (Conv2D)	(None, 128, 128, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 128, 128, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 128, 128, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 130, 130, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 64, 64, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 64, 64, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalization)	(None, 64, 64, 64)	256	conv2_block1_1_c
conv2_block1_1_relu (Activation)	(None, 64, 64, 64)	0	conv2_block1_1_t

conv2_block1_2_conv (Conv2D)	(None, 64, 64, 64)	36,928	conv2_block1_1_k conv2_block1_1_bn (BatchNormalization)
conv2_block1_2_bn (BatchNormalization)	(None, 64, 64, 64)	256	conv2_block1_2_k conv2_block1_2_relu (Activation)
conv2_block1_2_relu (Activation)	(None, 64, 64, 64)	0	conv2_block1_2_k conv2_block1_2_bn (BatchNormalization)
conv2_block1_0_conv (Conv2D)	(None, 64, 64, 256)	16,640	pool1_pool[0][0] conv2_block1_0_k conv2_block1_0_bn (BatchNormalization)
conv2_block1_3_conv (Conv2D)	(None, 64, 64, 256)	16,640	conv2_block1_2_k conv2_block1_3_k conv2_block1_3_bn (BatchNormalization)
conv2_block1_0_bn (BatchNormalization)	(None, 64, 64, 256)	1,024	conv2_block1_0_k conv2_block1_0_bn (BatchNormalization)
conv2_block1_3_bn (BatchNormalization)	(None, 64, 64, 256)	1,024	conv2_block1_3_k conv2_block1_3_add (Add)
conv2_block1_add (Add)	(None, 64, 64, 256)	0	conv2_block1_0_k conv2_block1_3_k conv2_block1_3_out (Activation)
conv2_block1_out (Activation)	(None, 64, 64, 256)	0	conv2_block1_add (Add)
conv2_block2_1_conv (Conv2D)	(None, 64, 64, 64)	16,448	conv2_block1_out (Activation)
conv2_block2_1_bn (BatchNormalization)	(None, 64, 64, 64)	256	conv2_block2_1_k conv2_block2_1_relu (Activation)
conv2_block2_1_relu (Activation)	(None, 64, 64, 64)	0	conv2_block2_1_k conv2_block2_2_conv (Conv2D)
conv2_block2_2_conv (Conv2D)	(None, 64, 64, 64)	36,928	conv2_block2_1_k conv2_block2_2_k conv2_block2_2_bn (BatchNormalization)
conv2_block2_2_bn (BatchNormalization)	(None, 64, 64, 64)	256	conv2_block2_2_k conv2_block2_2_relu (Activation)
conv2_block2_2_relu (Activation)	(None, 64, 64, 64)	0	conv2_block2_2_k conv2_block2_3_conv (Conv2D)
conv2_block2_3_conv (Conv2D)	(None, 64, 64, 256)	16,640	conv2_block2_2_k conv2_block2_3_k conv2_block2_3_bn (BatchNormalization)
conv2_block2_3_bn (BatchNormalization)	(None, 64, 64, 256)	1,024	conv2_block2_3_k conv2_block2_3_add (Add)
conv2_block2_add (Add)	(None, 64, 64, 256)	0	conv2_block1_out (Activation) conv2_block2_3_k

conv2_block2_out (Activation)	(None, 64, 64, 256)	0	conv2_block2_add
conv2_block3_1_conv (Conv2D)	(None, 64, 64, 64)	16,448	conv2_block2_out
conv2_block3_1_bn (BatchNormalization...)	(None, 64, 64, 64)	256	conv2_block3_1_c
conv2_block3_1_relu (Activation)	(None, 64, 64, 64)	0	conv2_block3_1_t
conv2_block3_2_conv (Conv2D)	(None, 64, 64, 64)	36,928	conv2_block3_1_i
conv2_block3_2_bn (BatchNormalizatio...)	(None, 64, 64, 64)	256	conv2_block3_2_c
conv2_block3_2_relu (Activation)	(None, 64, 64, 64)	0	conv2_block3_2_t
conv2_block3_3_conv (Conv2D)	(None, 64, 64, 256)	16,640	conv2_block3_2_i
conv2_block3_3_bn (BatchNormalizatio...)	(None, 64, 64, 256)	1,024	conv2_block3_3_c
conv2_block3_add (Add)	(None, 64, 64, 256)	0	conv2_block2_out conv2_block3_3_t
conv2_block3_out (Activation)	(None, 64, 64, 256)	0	conv2_block3_add
conv3_block1_1_conv (Conv2D)	(None, 32, 32, 128)	32,896	conv2_block3_out
conv3_block1_1_bn (BatchNormalizatio...)	(None, 32, 32, 128)	512	conv3_block1_1_c
conv3_block1_1_relu (Activation)	(None, 32, 32, 128)	0	conv3_block1_1_t
conv3_block1_2_conv (Conv2D)	(None, 32, 32, 128)	147,584	conv3_block1_1_i
conv3_block1_2_bn (BatchNormalizatio...)	(None, 32, 32, 128)	512	conv3_block1_2_c
conv3_block1_2_relu (Activation)	(None, 32, 32, 128)	0	conv3_block1_2_t
conv3_block1_0_conv	(None, 32, 32,	131,584	conv2_block3_out

(Conv2D)	512)		
conv3_block1_3_conv (Conv2D)	(None, 32, 32, 512)	66,048	conv3_block1_2_t
conv3_block1_0_bn (BatchNormalizatio...	(None, 32, 32, 512)	2,048	conv3_block1_0_c
conv3_block1_3_bn (BatchNormalizatio...	(None, 32, 32, 512)	2,048	conv3_block1_3_c
conv3_block1_add (Add)	(None, 32, 32, 512)	0	conv3_block1_0_t conv3_block1_3_t
conv3_block1_out (Activation)	(None, 32, 32, 512)	0	conv3_block1_adc
conv3_block2_1_conv (Conv2D)	(None, 32, 32, 128)	65,664	conv3_block1_out
conv3_block2_1_bn (BatchNormalizatio...	(None, 32, 32, 128)	512	conv3_block2_1_c
conv3_block2_1_relu (Activation)	(None, 32, 32, 128)	0	conv3_block2_1_t
conv3_block2_2_conv (Conv2D)	(None, 32, 32, 128)	147,584	conv3_block2_1_t
conv3_block2_2_bn (BatchNormalizatio...	(None, 32, 32, 128)	512	conv3_block2_2_c
conv3_block2_2_relu (Activation)	(None, 32, 32, 128)	0	conv3_block2_2_t
conv3_block2_3_conv (Conv2D)	(None, 32, 32, 512)	66,048	conv3_block2_2_t
conv3_block2_3_bn (BatchNormalizatio...	(None, 32, 32, 512)	2,048	conv3_block2_3_c
conv3_block2_add (Add)	(None, 32, 32, 512)	0	conv3_block1_out conv3_block2_3_t
conv3_block2_out (Activation)	(None, 32, 32, 512)	0	conv3_block2_adc
conv3_block3_1_conv (Conv2D)	(None, 32, 32, 128)	65,664	conv3_block2_out
conv3_block3_1_bn (BatchNormalizatio...	(None, 32, 32, 128)	512	conv3_block3_1_c

conv3_block3_1_relu (Activation)	(None, 32, 32, 128)	0	conv3_block3_1_t
conv3_block3_2_conv (Conv2D)	(None, 32, 32, 128)	147,584	conv3_block3_1_t
conv3_block3_2_bn (BatchNormalizatio...)	(None, 32, 32, 128)	512	conv3_block3_2_c
conv3_block3_2_relu (Activation)	(None, 32, 32, 128)	0	conv3_block3_2_t
conv3_block3_3_conv (Conv2D)	(None, 32, 32, 512)	66,048	conv3_block3_2_t
conv3_block3_3_bn (BatchNormalizatio...)	(None, 32, 32, 512)	2,048	conv3_block3_3_c
conv3_block3_add (Add)	(None, 32, 32, 512)	0	conv3_block2_out conv3_block3_3_t
conv3_block3_out (Activation)	(None, 32, 32, 512)	0	conv3_block3_adc
conv3_block4_1_conv (Conv2D)	(None, 32, 32, 128)	65,664	conv3_block3_out
conv3_block4_1_bn (BatchNormalizatio...)	(None, 32, 32, 128)	512	conv3_block4_1_c
conv3_block4_1_relu (Activation)	(None, 32, 32, 128)	0	conv3_block4_1_t
conv3_block4_2_conv (Conv2D)	(None, 32, 32, 128)	147,584	conv3_block4_1_t
conv3_block4_2_bn (BatchNormalizatio...)	(None, 32, 32, 128)	512	conv3_block4_2_c
conv3_block4_2_relu (Activation)	(None, 32, 32, 128)	0	conv3_block4_2_t
conv3_block4_3_conv (Conv2D)	(None, 32, 32, 512)	66,048	conv3_block4_2_t
conv3_block4_3_bn (BatchNormalizatio...)	(None, 32, 32, 512)	2,048	conv3_block4_3_c
conv3_block4_add (Add)	(None, 32, 32, 512)	0	conv3_block3_out conv3_block4_3_t
conv3_block4_out (Activation)	(None, 32, 32, 512)	0	conv3_block4_adc

conv4_block1_1_conv (Conv2D)	(None, 16, 16, 256)	131,328	conv3_block4_out
conv4_block1_1_bn (BatchNormalization...)	(None, 16, 16, 256)	1,024	conv4_block1_1_c
conv4_block1_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block1_1_t
conv4_block1_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block1_1_i
conv4_block1_2_bn (BatchNormalization...)	(None, 16, 16, 256)	1,024	conv4_block1_2_c
conv4_block1_2_relu (Activation)	(None, 16, 16, 256)	0	conv4_block1_2_t
conv4_block1_0_conv (Conv2D)	(None, 16, 16, 1024)	525,312	conv3_block4_out
conv4_block1_3_conv (Conv2D)	(None, 16, 16, 1024)	263,168	conv4_block1_2_i
conv4_block1_0_bn (BatchNormalization...)	(None, 16, 16, 1024)	4,096	conv4_block1_0_c
conv4_block1_3_bn (BatchNormalization...)	(None, 16, 16, 1024)	4,096	conv4_block1_3_c
conv4_block1_add (Add)	(None, 16, 16, 1024)	0	conv4_block1_0_t conv4_block1_3_t
conv4_block1_out (Activation)	(None, 16, 16, 1024)	0	conv4_block1_add
conv4_block2_1_conv (Conv2D)	(None, 16, 16, 256)	262,400	conv4_block1_out
conv4_block2_1_bn (BatchNormalization...)	(None, 16, 16, 256)	1,024	conv4_block2_1_c
conv4_block2_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block2_1_t
conv4_block2_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block2_1_i
conv4_block2_2_bn (BatchNormalization...)	(None, 16, 16, 256)	1,024	conv4_block2_2_c
conv4_block2_2_relu	(None, 16, 16,	0	conv4_block2_2_t

(Activation)	256)		
conv4_block2_3_conv (Conv2D)	(None, 16, 16, 1024)	263,168	conv4_block2_2_i
conv4_block2_3_bn (BatchNormalizatio...)	(None, 16, 16, 1024)	4,096	conv4_block2_3_c
conv4_block2_add (Add)	(None, 16, 16, 1024)	0	conv4_block1_out conv4_block2_3_k
conv4_block2_out (Activation)	(None, 16, 16, 1024)	0	conv4_block2_add
conv4_block3_1_conv (Conv2D)	(None, 16, 16, 256)	262,400	conv4_block2_out
conv4_block3_1_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block3_1_c
conv4_block3_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block3_1_k
conv4_block3_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block3_1_i
conv4_block3_2_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block3_2_c
conv4_block3_2_relu (Activation)	(None, 16, 16, 256)	0	conv4_block3_2_k
conv4_block3_3_conv (Conv2D)	(None, 16, 16, 1024)	263,168	conv4_block3_2_i
conv4_block3_3_bn (BatchNormalizatio...)	(None, 16, 16, 1024)	4,096	conv4_block3_3_c
conv4_block3_add (Add)	(None, 16, 16, 1024)	0	conv4_block2_out conv4_block3_3_k
conv4_block3_out (Activation)	(None, 16, 16, 1024)	0	conv4_block3_add
conv4_block4_1_conv (Conv2D)	(None, 16, 16, 256)	262,400	conv4_block3_out
conv4_block4_1_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block4_1_c
conv4_block4_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block4_1_k

conv4_block4_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block4_1_t
conv4_block4_2_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block4_2_c
conv4_block4_2_relu (Activation)	(None, 16, 16, 256)	0	conv4_block4_2_t
conv4_block4_3_conv (Conv2D)	(None, 16, 16, 1024)	263,168	conv4_block4_2_r
conv4_block4_3_bn (BatchNormalizatio...)	(None, 16, 16, 1024)	4,096	conv4_block4_3_c
conv4_block4_add (Add)	(None, 16, 16, 1024)	0	conv4_block3_out conv4_block4_3_t
conv4_block4_out (Activation)	(None, 16, 16, 1024)	0	conv4_block4_adv
conv4_block5_1_conv (Conv2D)	(None, 16, 16, 256)	262,400	conv4_block4_out conv4_block5_1_r
conv4_block5_1_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block5_1_c
conv4_block5_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block5_1_t
conv4_block5_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block5_1_r conv4_block5_2_t
conv4_block5_2_bn (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv4_block5_2_c
conv4_block5_2_relu (Activation)	(None, 16, 16, 256)	0	conv4_block5_2_t
conv4_block5_3_conv (Conv2D)	(None, 16, 16, 1024)	263,168	conv4_block5_2_r conv4_block5_3_t
conv4_block5_3_bn (BatchNormalizatio...)	(None, 16, 16, 1024)	4,096	conv4_block5_3_c
conv4_block5_add (Add)	(None, 16, 16, 1024)	0	conv4_block4_out conv4_block5_3_t
conv4_block5_out (Activation)	(None, 16, 16, 1024)	0	conv4_block5_adv conv4_block5_t
conv4_block6_1_conv (Conv2D)	(None, 16, 16, 256)	262,400	conv4_block5_out conv4_block6_1_t

conv4_block6_1_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv4_block6_1_t
conv4_block6_1_relu (Activation)	(None, 16, 16, 256)	0	conv4_block6_1_t
conv4_block6_2_conv (Conv2D)	(None, 16, 16, 256)	590,080	conv4_block6_1_t
conv4_block6_2_bn (BatchNormalization)	(None, 16, 16, 256)	1,024	conv4_block6_2_t
conv4_block6_2_relu (Activation)	(None, 16, 16, 256)	0	conv4_block6_2_t
average_pooling2d (AveragePooling2D)	(None, 1, 1, 256)	0	conv4_block6_2_t
conv2d_49 (Conv2D)	(None, 1, 1, 256)	65,536	average_pooling2
batch_normalizatio... (BatchNormalizatio...)	(None, 1, 1, 256)	1,024	conv2d_49[0][0]
conv2d_50 (Conv2D)	(None, 16, 16, 256)	65,536	conv4_block6_2_t
conv2d_51 (Conv2D)	(None, 16, 16, 256)	65,536	conv4_block6_2_t
conv2d_52 (Conv2D)	(None, 16, 16, 256)	65,536	conv4_block6_2_t
conv2d_53 (Conv2D)	(None, 16, 16, 256)	65,536	conv4_block6_2_t
activation_8 (Activation)	(None, 1, 1, 256)	0	batch_normalizati
batch_normalizatio... (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv2d_50[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv2d_51[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv2d_52[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv2d_53[0][0]
up_sampling2d (UpSampling2D)	(None, 16, 16, 256)	0	activation_8[0][0]

activation_9 (Activation)	(None, 16, 16, 256)	0	batch_normalization_9[0]
activation_10 (Activation)	(None, 16, 16, 256)	0	batch_normalization_10[0]
activation_11 (Activation)	(None, 16, 16, 256)	0	batch_normalization_11[0]
activation_12 (Activation)	(None, 16, 16, 256)	0	batch_normalization_12[0]
concatenate_10 (Concatenate)	(None, 16, 16, 1280)	0	up_sampling2d[0] activation_9[0] activation_10[0] activation_11[0] activation_12[0]
conv2d_54 (Conv2D)	(None, 16, 16, 256)	327,680	concatenate_10[0]
batch_normalizatio... (BatchNormalizatio...)	(None, 16, 16, 256)	1,024	conv2d_54[0][0]
conv2d_55 (Conv2D)	(None, 64, 64, 48)	3,072	conv2_block3_2[0]
activation_13 (Activation)	(None, 16, 16, 256)	0	batch_normalization_13[0]
batch_normalizatio... (BatchNormalizatio...)	(None, 64, 64, 48)	192	conv2d_55[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 64, 64, 256)	0	activation_13[0]
activation_14 (Activation)	(None, 64, 64, 48)	0	batch_normalization_14[0]
concatenate_11 (Concatenate)	(None, 64, 64, 304)	0	up_sampling2d_1[0] activation_14[0]
conv2d_56 (Conv2D)	(None, 64, 64, 256)	700,416	concatenate_11[0]
batch_normalizatio... (BatchNormalizatio...)	(None, 64, 64, 256)	1,024	conv2d_56[0][0]
activation_15 (Activation)	(None, 64, 64, 256)	0	batch_normalization_15[0]
conv2d_57 (Conv2D)	(None, 64, 64, 256)	589,824	activation_15[0]

	256)		
batch_normalization... (BatchNormalizatio...)	(None, 64, 64, 256)	1,024	conv2d_57[0][0]
activation_16 (Activation)	(None, 64, 64, 256)	0	batch_normalizati
up_sampling2d_2 (UpSampling2D)	(None, 256, 256, 256)	0	activation_16[0]
output_layer (Conv2D)	(None, 256, 256, 13)	3,341	up_sampling2d_2[0]
activation_17 (Activation)	(None, 256, 256, 13)	0	output_layer[0]

Total params: 10,282,317 (39.22 MB)

Trainable params: 10,249,581 (39.10 MB)

Non-trainable params: 32,736 (127.88 KB)

In [24]:

```
start_time = time.time()

early_stop = EarlyStopping(monitor='val_accuracy', patience=5, restore_
best_weights=True)
reduce_lr = ReduceLROnPlateau(monitor='val_accuracy', factor=1e-1, patie
nce=10, verbose=1, min_lr = 1e-6)

deeplab_result1 = deeplab_model1.fit(train_dataset,
                                      validation_data = val_dataset,
                                      epochs = EPOCHS,
                                      callbacks=[early_stop, reduce_lr],
                                      batch_size = BATCH_SIZE)
end_time = time.time()
```

Epoch 1/15

110/110  **248s** 1s/step - accuracy: 0.8340 - loss:
0.5719 - val_accuracy: 0.3233 - val_loss: 2.0047 - learning_rate: 0.001
0

Epoch 2/15

110/110  **101s** 842ms/step - accuracy: 0.9427 - los
s: 0.1571 - val_accuracy: 0.4221 - val_loss: 2.1206 - learning_rate: 0.
0010

Epoch 3/15

110/110  **100s** 838ms/step - accuracy: 0.9516 - los
s: 0.1269 - val_accuracy: 0.4205 - val_loss: 2.4646 - learning_rate: 0.
0010

Epoch 4/15

110/110  **100s** 837ms/step - accuracy: 0.9558 - los
s: 0.1135 - val_accuracy: 0.4405 - val_loss: 2.1882 - learning_rate: 0.
0010

Epoch 5/15

110/110  **100s** 840ms/step - accuracy: 0.9576 - los
s: 0.1077 - val_accuracy: 0.5647 - val_loss: 1.4782 - learning_rate: 0.
0010

Epoch 6/15

110/110 ————— **100s** 840ms/step - accuracy: 0.9595 - loss: 0.1018 - val_accuracy: 0.6184 - val_loss: 1.4269 - learning_rate: 0.0010
Epoch 7/15
110/110 ————— **101s** 840ms/step - accuracy: 0.9611 - loss: 0.0972 - val_accuracy: 0.7973 - val_loss: 0.7535 - learning_rate: 0.0010
Epoch 8/15
110/110 ————— **101s** 843ms/step - accuracy: 0.9622 - loss: 0.0937 - val_accuracy: 0.8655 - val_loss: 0.4462 - learning_rate: 0.0010
Epoch 9/15
110/110 ————— **100s** 841ms/step - accuracy: 0.9624 - loss: 0.0928 - val_accuracy: 0.9152 - val_loss: 0.2511 - learning_rate: 0.0010
Epoch 10/15
110/110 ————— **101s** 842ms/step - accuracy: 0.9631 - loss: 0.0906 - val_accuracy: 0.9358 - val_loss: 0.1827 - learning_rate: 0.0010
Epoch 11/15
110/110 ————— **100s** 839ms/step - accuracy: 0.9636 - loss: 0.0892 - val_accuracy: 0.9519 - val_loss: 0.1265 - learning_rate: 0.0010
Epoch 12/15
110/110 ————— **100s** 839ms/step - accuracy: 0.9640 - loss: 0.0878 - val_accuracy: 0.9587 - val_loss: 0.1035 - learning_rate: 0.0010
Epoch 13/15
110/110 ————— **100s** 836ms/step - accuracy: 0.9644 - loss: 0.0866 - val_accuracy: 0.9578 - val_loss: 0.1059 - learning_rate: 0.0010
Epoch 14/15
110/110 ————— **100s** 839ms/step - accuracy: 0.9645 - loss: 0.0861 - val_accuracy: 0.9407 - val_loss: 0.2083 - learning_rate: 0.0010
Epoch 15/15
110/110 ————— **101s** 842ms/step - accuracy: 0.9607 - loss:

```
s: 0.0993 - val_accuracy: 0.8211 - val_loss: 0.7244 - learning_rate: 0.  
0010
```

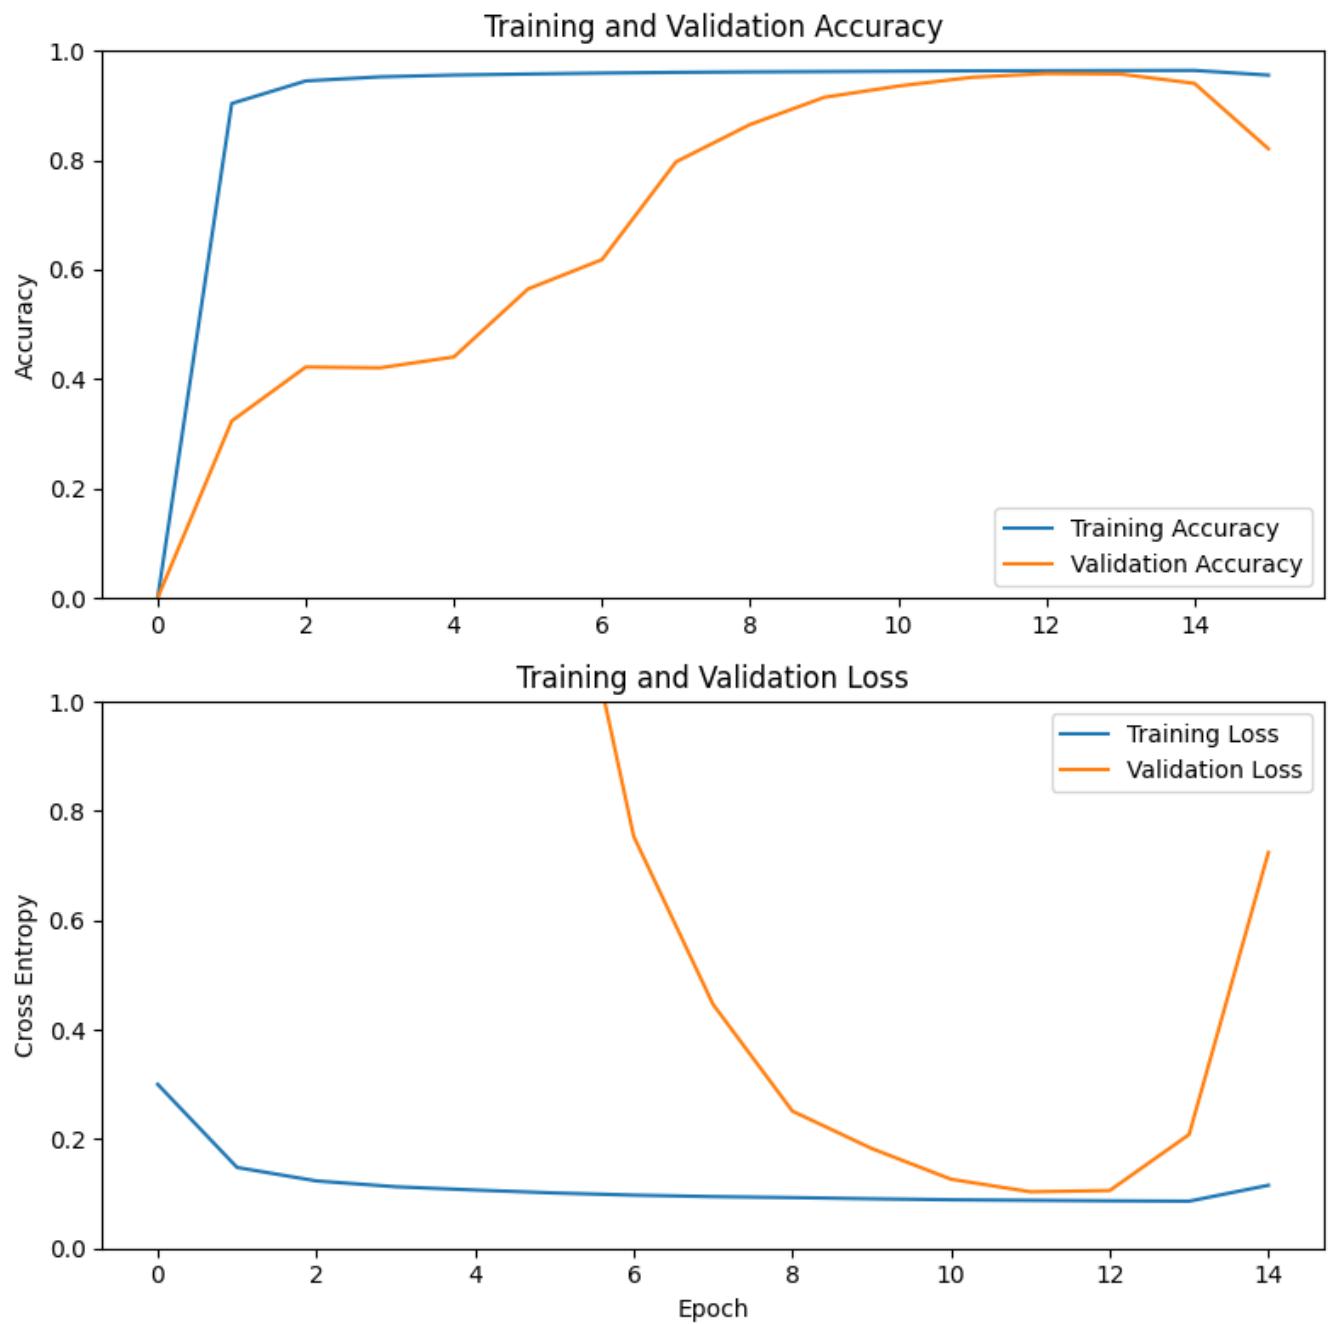
In [25]:

```
training_time = end_time - start_time  
print("-----")  
print("Training time for Deeplab:", training_time, "seconds")
```

```
-----  
Training time for Deeplab: 1653.4657363891602 seconds
```

In [26]:

```
plot_training_history(deeplab_result1)
```



4. Evaluation

- ### Accuracy

In [27]:

```
def model_accuracy(model):
    train_loss, train_accuracy = model.evaluate(train_dataset, batch_size = BATCH_SIZE)
    validation_loss, validation_accuracy = model.evaluate(val_dataset, batch_size = BATCH_SIZE)
    test_loss, test_accuracy = model.evaluate(test_dataset, batch_size = BATCH_SIZE)
    print(f'Model Accuracy on the Training Dataset: {round(train_accuracy * 100, 2)}%')
    print(f'Model Accuracy on the Validation Dataset: {round(validation_accuracy * 100, 2)}%')
    print(f'Model Accuracy on the Test Dataset: {round(test_accuracy * 100, 2)}%)
```

In [28]:

```
model_accuracy(fcn_model1)
```

110/110 ————— **39s** 286ms/step - accuracy: 0.9089 - loss: 0.3347

24/24 ————— **41s** 220ms/step - accuracy: 0.9044 - loss: 0.3474

24/24 ————— **43s** 118ms/step - accuracy: 0.9088 - loss: 0.3353

Model Accuracy on the Training Dataset: 90.92%

Model Accuracy on the Validation Dataset: 90.58%

Model Accuracy on the Test Dataset: 90.85%

In [29]:

```
model_accuracy(unet_model1)
```

110/110 ————— **40s** 288ms/step - accuracy: 0.9794 - loss: 0.0610
24/24 ————— **41s** 213ms/step - accuracy: 0.9777 - loss: 0.0673
24/24 ————— **42s** 104ms/step - accuracy: 0.9783 - loss: 0.0656
Model Accuracy on the Training Dataset: 97.95%
Model Accuracy on the Validation Dataset: 97.73%
Model Accuracy on the Test Dataset: 97.83%

In [30]:

```
model_accuracy(deeplab_model1)
```

110/110 ————— **43s** 322ms/step - accuracy: 0.9617 - loss: 0.0934
24/24 ————— **42s** 263ms/step - accuracy: 0.9594 - loss: 0.1015
24/24 ————— **44s** 202ms/step - accuracy: 0.9602 - loss: 0.0991
Model Accuracy on the Training Dataset: 96.16%
Model Accuracy on the Validation Dataset: 95.87%
Model Accuracy on the Test Dataset: 96.0%

- ### Intersection-over-Union (IoU)

In [31]:

```
import numpy as np
import tensorflow as tf
import pandas as pd
import numpy as np
```

```
def evaluate_metrics(dataset, model, n_classes=13):
    # create mask
    true_masks, predicted_masks = [], []
    for images, masks in dataset:
        pred_masks = model.predict(images)
        pred_masks = tf.expand_dims(tf.argmax(pred_masks, axis=-1), axis=-1)
        true_masks.append(masks)
        predicted_masks.append(pred_masks)
    true_masks = np.array(true_masks)
    predicted_masks = np.array(predicted_masks)

    # initialize
    class_wise_metrics = {
        "TP": np.zeros(n_classes), # True Positives
        "FP": np.zeros(n_classes), # False Positives
        "FN": np.zeros(n_classes) # False Negatives
    }

    # calculate metrics per class
    for c_id in range(n_classes):
        true_positive = (predicted_masks == c_id) & (true_masks == c_id)
        false_positive = (predicted_masks == c_id) & (true_masks != c_id)
        false_negative = (predicted_masks != c_id) & (true_masks == c_id)

        class_wise_metrics["TP"][c_id] += np.sum(true_positive)
        class_wise_metrics["FP"][c_id] += np.sum(false_positive)
        class_wise_metrics["FN"][c_id] += np.sum(false_negative)

    # class-wise metrics
    recall = np.round(class_wise_metrics["TP"] / (class_wise_metrics["TP"] + class_wise_metrics["FN"] + 1e-6), 2)
    precision = np.round(class_wise_metrics["TP"] / (class_wise_metrics["TP"] + class_wise_metrics["FP"] + 1e-6), 2)
```

```
[ "TP" ] + class_wise_metrics[ "FP" ] + 1e-6), 2)
    iou = np.round(class_wise_metrics[ "TP" ] / (class_wise_metrics[ "TP" ]
+ class_wise_metrics[ "FP" ] + class_wise_metrics[ "FN" ] + 1e-6), 2)

    # overall metrics
    overall_recall = np.mean(recall)
    overall_precision = np.mean(precision)
    overall_iou = np.mean(iou)

    # package evaluations
    evaluations = {
        "class_wise": {
            "Recall": recall.tolist(),
            "Precision": precision.tolist(),
            "IoU": iou.tolist()
        },
        "overall": {
            "Recall": round(overall_recall, 2),
            "Precision": round(overall_precision, 2),
            "IoU": round(overall_iou, 2)
        }
    }

    return evaluations

def show_evaluations(evaluations):
    class_wise_data = evaluations[ 'class_wise' ]
    class_wise_df = pd.DataFrame(class_wise_data)
    class_wise_df.index = [f"Class {i+1}" for i in range(len(class_wise_data[ 'Recall' ]))]

    overall_data = evaluations[ 'overall' ]
    overall_df = pd.DataFrame(overall_data, index=[ 'All Classes' ])

    # Combine
    combined_df = pd.concat([overall_df, class_wise_df]).reset_index()
```

```
combined_df.rename(columns={'index': 'Class'}, inplace=True)

return combined_df
```

fcn

In [32]:

```
evaluations = evaluate_metrics(train_dataset, fcn_model1)
show_evaluations(evaluations)
```

1/1	1s	709ms/step
1/1	0s	61ms/step
1/1	0s	55ms/step
1/1	0s	56ms/step
1/1	0s	36ms/step
1/1	0s	34ms/step
1/1	0s	37ms/step
1/1	0s	38ms/step
1/1	0s	43ms/step
1/1	0s	41ms/step
1/1	0s	42ms/step
1/1	0s	36ms/step
1/1	0s	42ms/step
1/1	0s	52ms/step
1/1	0s	35ms/step
1/1	0s	41ms/step
1/1	0s	35ms/step
1/1	0s	37ms/step
1/1	0s	35ms/step
1/1	0s	41ms/step
1/1	0s	34ms/step
1/1	0s	38ms/step
1/1	0s	42ms/step
1/1	0s	45ms/step

1/1 0s 36ms/step
1/1 0s 51ms/step
1/1 0s 35ms/step
1/1 0s 38ms/step
1/1 0s 47ms/step
1/1 0s 40ms/step
1/1 0s 55ms/step
1/1 0s 36ms/step
1/1 0s 41ms/step
1/1 0s 37ms/step
1/1 0s 40ms/step
1/1 0s 50ms/step
1/1 0s 34ms/step
1/1 0s 41ms/step
1/1 0s 64ms/step
1/1 0s 41ms/step
1/1 0s 45ms/step
1/1 0s 39ms/step
1/1 0s 40ms/step
1/1 0s 45ms/step
1/1 0s 35ms/step
1/1 0s 38ms/step
1/1 0s 37ms/step
1/1 0s 40ms/step
1/1 0s 43ms/step
1/1 0s 42ms/step
1/1 0s 45ms/step
1/1 0s 34ms/step
1/1 0s 36ms/step
1/1 0s 35ms/step
1/1 0s 39ms/step
1/1 0s 37ms/step
1/1 0s 40ms/step
1/1 0s 38ms/step
1/1 0s 45ms/step
1/1 0s 44ms/step
1/1 0s 47ms/step

1/1 0s 44ms/step
1/1 0s 37ms/step
1/1 0s 52ms/step
1/1 0s 40ms/step
1/1 0s 35ms/step
1/1 0s 37ms/step
1/1 0s 39ms/step
1/1 0s 40ms/step
1/1 0s 35ms/step
1/1 0s 30ms/step
1/1 0s 42ms/step
1/1 0s 36ms/step
1/1 0s 47ms/step
1/1 0s 38ms/step
1/1 0s 36ms/step
1/1 0s 39ms/step
1/1 0s 49ms/step
1/1 0s 37ms/step
1/1 0s 43ms/step
1/1 0s 33ms/step
1/1 0s 38ms/step
1/1 0s 40ms/step
1/1 0s 57ms/step
1/1 0s 71ms/step
1/1 0s 39ms/step
1/1 0s 40ms/step
1/1 0s 38ms/step
1/1 0s 45ms/step
1/1 0s 35ms/step
1/1 0s 53ms/step
1/1 0s 38ms/step
1/1 0s 37ms/step
1/1 0s 36ms/step
1/1 0s 39ms/step
1/1 0s 35ms/step
1/1 0s 38ms/step
1/1 0s 43ms/step

```
1/1 ━━━━━━ 0s 35ms/step
1/1 ━━━━━━ 0s 41ms/step
1/1 ━━━━━━ 0s 54ms/step
1/1 ━━━━━━ 0s 38ms/step
1/1 ━━━━━━ 0s 35ms/step
1/1 ━━━━━━ 0s 40ms/step
1/1 ━━━━━━ 0s 36ms/step
1/1 ━━━━━━ 0s 40ms/step
1/1 ━━━━━━ 0s 35ms/step
1/1 ━━━━━━ 0s 39ms/step
1/1 ━━━━━━ 0s 38ms/step
1/1 ━━━━━━ 0s 499ms/step
```

Out[32]:

	Class	Recall	Precision	IoU
0	All Classes	0.61	0.75	0.55
1	Class 1	0.97	0.98	0.96
2	Class 2	0.83	0.77	0.66
3	Class 3	0.46	0.73	0.40
4	Class 4	0.29	0.59	0.24
5	Class 5	0.00	0.00	0.00
6	Class 6	0.33	0.60	0.27
7	Class 7	0.89	0.98	0.88
8	Class 8	0.93	0.91	0.86
9	Class 9	0.79	0.82	0.68
10	Class 10	0.92	0.87	0.81
11	Class 11	0.96	0.94	0.90
12	Class 12	0.37	0.72	0.32
13	Class 13	0.21	0.78	0.20

In [33]:

```
evaluations = evaluate_metrics(val_dataset, fcn_model1)
show_evaluations(evaluations)
```

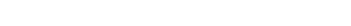
1/1  **0s** 59ms/step
1/1  **0s** 48ms/step
1/1  **0s** 52ms/step
1/1  **0s** 39ms/step
1/1  **0s** 35ms/step
1/1  **0s** 38ms/step
1/1  **0s** 39ms/step
1/1  **0s** 36ms/step
1/1  **0s** 35ms/step
1/1  **0s** 39ms/step
1/1  **0s** 36ms/step
1/1  **0s** 35ms/step
1/1  **0s** 43ms/step
1/1  **0s** 36ms/step
1/1  **0s** 28ms/step
1/1  **0s** 32ms/step
1/1  **0s** 28ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 276ms/step

Out[33]:

	Class	Recall	Precision	IoU
0	All Classes	0.61	0.74	0.54
1	Class 1	0.97	0.98	0.95
2	Class 2	0.81	0.76	0.65
3	Class 3	0.48	0.74	0.41
4	Class 4	0.26	0.56	0.22
5	Class 5	0.00	0.00	0.00
6	Class 6	0.33	0.60	0.27
7	Class 7	0.88	0.98	0.87
8	Class 8	0.93	0.91	0.85
9	Class 9	0.79	0.81	0.66
10	Class 10	0.92	0.87	0.81
11	Class 11	0.96	0.94	0.90
12	Class 12	0.36	0.72	0.31
13	Class 13	0.19	0.73	0.18

In [34]:

```
evaluations = evaluate_metrics(test_dataset, fcn_model1)
show_evaluations(evaluations)
```

1/1  **0s** 29ms/step
1/1  **0s** 43ms/step
1/1  **0s** 29ms/step
1/1  **0s** 34ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 41ms/step
1/1  **0s** 29ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 28ms/step
1/1  **0s** 28ms/step
1/1  **0s** 28ms/step
1/1  **0s** 28ms/step
1/1  **0s** 29ms/step
1/1  **0s** 33ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 28ms/step
1/1  **0s** 29ms/step
1/1  **0s** 24ms/step

Out[34]:

	Class	Recall	Precision	IoU
0	All Classes	0.61	0.74	0.55
1	Class 1	0.97	0.98	0.96
2	Class 2	0.82	0.78	0.67
3	Class 3	0.47	0.74	0.40
4	Class 4	0.29	0.59	0.24
5	Class 5	0.00	0.00	0.00
6	Class 6	0.33	0.59	0.27
7	Class 7	0.90	0.98	0.88
8	Class 8	0.93	0.91	0.85
9	Class 9	0.79	0.82	0.68
10	Class 10	0.92	0.87	0.81
11	Class 11	0.97	0.93	0.90
12	Class 12	0.35	0.71	0.31
13	Class 13	0.22	0.71	0.20

unet

In [35]:

```
evaluations = evaluate_metrics(train_dataset, unet_model1)
show_evaluations(evaluations)
```

1/1 ━━━━━━ **2s** 2s/step
1/1 ━━━━━━ **0s** 53ms/step
1/1 ━━━━━━ **0s** 42ms/step
1/1 ━━━━━━ **0s** 39ms/step
1/1 ━━━━━━ **0s** 38ms/step
1/1 ━━━━━━ **0s** 69ms/step
1/1 ━━━━━━ **0s** 40ms/step
1/1 ━━━━━━ **0s** 51ms/step
1/1 ━━━━━━ **0s** 43ms/step

1/1 0s 70ms/step
1/1 0s 38ms/step
1/1 0s 41ms/step
1/1 0s 70ms/step
1/1 0s 44ms/step
1/1 0s 40ms/step
1/1 0s 54ms/step
1/1 0s 42ms/step
1/1 0s 46ms/step
1/1 0s 37ms/step
1/1 0s 39ms/step
1/1 0s 49ms/step
1/1 0s 42ms/step
1/1 0s 50ms/step
1/1 0s 40ms/step
1/1 0s 41ms/step
1/1 0s 38ms/step
1/1 0s 51ms/step
1/1 0s 41ms/step
1/1 0s 37ms/step
1/1 0s 54ms/step
1/1 0s 53ms/step
1/1 0s 46ms/step
1/1 0s 60ms/step
1/1 0s 49ms/step
1/1 0s 49ms/step
1/1 0s 61ms/step
1/1 0s 45ms/step
1/1 0s 37ms/step
1/1 0s 43ms/step
1/1 0s 55ms/step
1/1 0s 46ms/step
1/1 0s 38ms/step
1/1 0s 43ms/step
1/1 0s 43ms/step
1/1 0s 45ms/step
1/1 0s 41ms/step

1/1 0s 47ms/step
1/1 0s 53ms/step
1/1 0s 54ms/step
1/1 0s 57ms/step
1/1 0s 48ms/step
1/1 0s 47ms/step
1/1 0s 41ms/step
1/1 0s 41ms/step
1/1 0s 38ms/step
1/1 0s 41ms/step
1/1 0s 48ms/step
1/1 0s 37ms/step
1/1 0s 37ms/step
1/1 0s 37ms/step
1/1 0s 39ms/step
1/1 0s 55ms/step
1/1 0s 37ms/step
1/1 0s 37ms/step
1/1 0s 39ms/step
1/1 0s 44ms/step
1/1 0s 47ms/step
1/1 0s 41ms/step
1/1 0s 38ms/step
1/1 0s 39ms/step
1/1 0s 43ms/step
1/1 0s 48ms/step
1/1 0s 41ms/step
1/1 0s 38ms/step
1/1 0s 52ms/step
1/1 0s 42ms/step
1/1 0s 37ms/step
1/1 0s 51ms/step
1/1 0s 37ms/step
1/1 0s 41ms/step
1/1 0s 40ms/step
1/1 0s 44ms/step
1/1 0s 65ms/step

1/1 0s 49ms/step
1/1 0s 57ms/step
1/1 0s 43ms/step
1/1 0s 41ms/step
1/1 0s 39ms/step
1/1 0s 39ms/step
1/1 0s 40ms/step
1/1 0s 41ms/step
1/1 0s 45ms/step
1/1 0s 50ms/step
1/1 0s 34ms/step
1/1 0s 41ms/step
1/1 0s 45ms/step
1/1 0s 56ms/step
1/1 0s 49ms/step
1/1 0s 51ms/step
1/1 0s 46ms/step
1/1 0s 37ms/step
1/1 0s 47ms/step
1/1 0s 43ms/step
1/1 0s 38ms/step
1/1 0s 63ms/step
1/1 0s 44ms/step
1/1 0s 45ms/step
1/1 0s 46ms/step
1/1 0s 30ms/step
1/1 1s 1s/step

Out[35]:

	Class	Recall	Precision	IoU
0	All Classes	0.82	0.87	0.78
1	Class 1	0.99	0.99	0.98
2	Class 2	0.98	0.95	0.94
3	Class 3	0.76	0.84	0.66
4	Class 4	0.74	0.91	0.69
5	Class 5	0.00	0.00	0.00
6	Class 6	0.66	0.90	0.61
7	Class 7	0.97	0.98	0.95
8	Class 8	1.00	1.00	0.99
9	Class 9	0.98	0.97	0.95
10	Class 10	0.96	0.96	0.92
11	Class 11	0.99	0.99	0.99
12	Class 12	0.87	0.92	0.80
13	Class 13	0.72	0.86	0.65

In [36]:

```
evaluations = evaluate_metrics(val_dataset, unet_model1)
show_evaluations(evaluations)
```

1/1  **0s** 35ms/step
1/1  **0s** 61ms/step
1/1  **0s** 38ms/step
1/1  **0s** 47ms/step
1/1  **0s** 44ms/step
1/1  **0s** 46ms/step
1/1  **0s** 45ms/step
1/1  **0s** 47ms/step
1/1  **0s** 45ms/step
1/1  **0s** 51ms/step
1/1  **0s** 41ms/step
1/1  **0s** 56ms/step
1/1  **0s** 38ms/step
1/1  **0s** 59ms/step
1/1  **0s** 50ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 29ms/step
1/1  **0s** 30ms/step
1/1  **1s** 749ms/step

Out[36]:

	Class	Recall	Precision	IoU
0	All Classes	0.80	0.86	0.76
1	Class 1	0.99	0.99	0.98
2	Class 2	0.98	0.94	0.92
3	Class 3	0.76	0.83	0.66
4	Class 4	0.68	0.89	0.62
5	Class 5	0.00	0.00	0.00
6	Class 6	0.62	0.89	0.57
7	Class 7	0.97	0.97	0.94
8	Class 8	0.99	0.99	0.99
9	Class 9	0.98	0.96	0.94
10	Class 10	0.96	0.96	0.92
11	Class 11	0.99	0.99	0.99
12	Class 12	0.85	0.90	0.78
13	Class 13	0.68	0.85	0.60

In [37]:

```
evaluations = evaluate_metrics(test_dataset, unet_model1)
show_evaluations(evaluations)
```

1/1  **0s** 32ms/step
1/1  **0s** 32ms/step
1/1  **0s** 51ms/step
1/1  **0s** 39ms/step
1/1  **0s** 30ms/step
1/1  **0s** 30ms/step
1/1  **0s** 31ms/step
1/1  **0s** 30ms/step
1/1  **0s** 30ms/step
1/1  **0s** 30ms/step
1/1  **0s** 33ms/step
1/1  **0s** 31ms/step
1/1  **0s** 30ms/step
1/1  **0s** 30ms/step
1/1  **0s** 32ms/step
1/1  **0s** 29ms/step
1/1  **0s** 31ms/step
1/1  **0s** 31ms/step
1/1  **0s** 30ms/step
1/1  **0s** 31ms/step
1/1  **0s** 30ms/step
1/1  **0s** 29ms/step
1/1  **0s** 27ms/step

Out[37]:

	Class	Recall	Precision	IoU
0	All Classes	0.81	0.86	0.77
1	Class 1	0.99	0.99	0.98
2	Class 2	0.98	0.95	0.93
3	Class 3	0.76	0.83	0.66
4	Class 4	0.72	0.90	0.67
5	Class 5	0.00	0.00	0.00
6	Class 6	0.62	0.89	0.58
7	Class 7	0.97	0.98	0.95
8	Class 8	0.99	1.00	0.99
9	Class 9	0.98	0.97	0.94
10	Class 10	0.96	0.96	0.92
11	Class 11	0.99	0.99	0.99
12	Class 12	0.85	0.90	0.78
13	Class 13	0.67	0.83	0.58

deeplab

In [38]:

```
evaluations = evaluate_metrics(train_dataset, deeplab_model1)
show_evaluations(evaluations)
```

1/1 ━━━━━━ **4s** 4s/step
1/1 ━━━━━━ **0s** 42ms/step
1/1 ━━━━━━ **0s** 47ms/step
1/1 ━━━━━━ **0s** 43ms/step
1/1 ━━━━━━ **0s** 49ms/step
1/1 ━━━━━━ **0s** 50ms/step
1/1 ━━━━━━ **0s** 46ms/step
1/1 ━━━━━━ **0s** 41ms/step
1/1 ━━━━━━ **0s** 50ms/step

1/1 0s 43ms/step
1/1 0s 64ms/step
1/1 0s 51ms/step
1/1 0s 44ms/step
1/1 0s 42ms/step
1/1 0s 44ms/step
1/1 0s 40ms/step
1/1 0s 43ms/step
1/1 0s 54ms/step
1/1 0s 46ms/step
1/1 0s 45ms/step
1/1 0s 46ms/step
1/1 0s 46ms/step
1/1 0s 42ms/step
1/1 0s 40ms/step
1/1 0s 45ms/step
1/1 0s 49ms/step
1/1 0s 46ms/step
1/1 0s 52ms/step
1/1 0s 52ms/step
1/1 0s 51ms/step
1/1 0s 42ms/step
1/1 0s 43ms/step
1/1 0s 46ms/step
1/1 0s 46ms/step
1/1 0s 49ms/step
1/1 0s 50ms/step
1/1 0s 49ms/step
1/1 0s 45ms/step
1/1 0s 45ms/step
1/1 0s 47ms/step
1/1 0s 51ms/step
1/1 0s 43ms/step
1/1 0s 49ms/step
1/1 0s 46ms/step
1/1 0s 49ms/step
1/1 0s 58ms/step

1/1 0s 42ms/step
1/1 0s 41ms/step
1/1 0s 40ms/step
1/1 0s 49ms/step
1/1 0s 42ms/step
1/1 0s 45ms/step
1/1 0s 42ms/step
1/1 0s 47ms/step
1/1 0s 42ms/step
1/1 0s 41ms/step
1/1 0s 53ms/step
1/1 0s 54ms/step
1/1 0s 59ms/step
1/1 0s 45ms/step
1/1 0s 40ms/step
1/1 0s 42ms/step
1/1 0s 50ms/step
1/1 0s 48ms/step
1/1 0s 49ms/step
1/1 0s 49ms/step
1/1 0s 47ms/step
1/1 0s 46ms/step
1/1 0s 71ms/step
1/1 0s 50ms/step
1/1 0s 42ms/step
1/1 0s 43ms/step
1/1 0s 46ms/step
1/1 0s 43ms/step
1/1 0s 45ms/step
1/1 0s 40ms/step
1/1 0s 43ms/step
1/1 0s 52ms/step
1/1 0s 43ms/step
1/1 0s 44ms/step
1/1 0s 43ms/step
1/1 0s 46ms/step
1/1 0s 41ms/step

1/1  **0s** 36ms/step
1/1  **0s** 43ms/step
1/1  **0s** 44ms/step
1/1  **0s** 39ms/step
1/1  **0s** 40ms/step
1/1  **0s** 44ms/step
1/1  **0s** 56ms/step
1/1  **0s** 42ms/step
1/1  **0s** 53ms/step
1/1  **0s** 45ms/step
1/1  **0s** 46ms/step
1/1  **0s** 41ms/step
1/1  **0s** 42ms/step
1/1  **0s** 57ms/step
1/1  **0s** 35ms/step
1/1  **0s** 40ms/step
1/1  **0s** 42ms/step
1/1  **0s** 43ms/step
1/1  **0s** 48ms/step
1/1  **0s** 50ms/step
1/1  **0s** 48ms/step
1/1  **0s** 42ms/step
1/1  **0s** 62ms/step
1/1  **0s** 48ms/step
1/1  **0s** 43ms/step
1/1  **0s** 44ms/step
1/1  **4s** 4s/step

Out[38]:

	Class	Recall	Precision	IoU
0	All Classes	0.86	0.84	0.76
1	Class 1	0.97	0.98	0.95
2	Class 2	0.96	0.97	0.93
3	Class 3	0.74	0.63	0.52
4	Class 4	0.85	0.81	0.71
5	Class 5	0.62	0.56	0.41
6	Class 6	0.73	0.68	0.54
7	Class 7	0.79	0.75	0.63
8	Class 8	0.98	0.99	0.97
9	Class 9	0.96	0.97	0.93
10	Class 10	0.93	0.90	0.84
11	Class 11	0.99	0.99	0.99
12	Class 12	0.93	0.85	0.80
13	Class 13	0.72	0.88	0.66

In [39]:

```
evaluations = evaluate_metrics(val_dataset, deeplab_model1)
show_evaluations(evaluations)
```

1/1  **0s** 55ms/step
1/1  **0s** 55ms/step
1/1  **0s** 42ms/step
1/1  **0s** 65ms/step
1/1  **0s** 46ms/step
1/1  **0s** 47ms/step
1/1  **0s** 51ms/step
1/1  **0s** 44ms/step
1/1  **0s** 46ms/step
1/1  **0s** 46ms/step
1/1  **0s** 72ms/step
1/1  **0s** 45ms/step
1/1  **0s** 43ms/step
1/1  **0s** 41ms/step
1/1  **0s** 48ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **2s** 2s/step

Out[39]:

	Class	Recall	Precision	IoU
0	All Classes	0.84	0.83	0.74
1	Class 1	0.97	0.98	0.95
2	Class 2	0.95	0.96	0.92
3	Class 3	0.74	0.62	0.51
4	Class 4	0.80	0.78	0.65
5	Class 5	0.53	0.56	0.37
6	Class 6	0.67	0.67	0.50
7	Class 7	0.79	0.75	0.62
8	Class 8	0.98	0.99	0.97
9	Class 9	0.95	0.97	0.92
10	Class 10	0.93	0.89	0.83
11	Class 11	0.99	0.99	0.99
12	Class 12	0.91	0.83	0.76
13	Class 13	0.66	0.86	0.60

In [40]:

```
evaluations = evaluate_metrics(test_dataset, deeplab_model1)
show_evaluations(evaluations)
```

1/1  **0s** 34ms/step
1/1  **0s** 35ms/step
1/1  **0s** 34ms/step
1/1  **0s** 37ms/step
1/1  **0s** 32ms/step
1/1  **0s** 32ms/step
1/1  **0s** 46ms/step
1/1  **0s** 40ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 46ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 32ms/step
1/1  **0s** 32ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 32ms/step
1/1  **0s** 33ms/step
1/1  **0s** 33ms/step
1/1  **0s** 27ms/step

Out[40]:

	Class	Recall	Precision	IoU
0	All Classes	0.81	0.84	0.73
1	Class 1	0.97	0.98	0.95
2	Class 2	0.96	0.97	0.93
3	Class 3	0.74	0.62	0.51
4	Class 4	0.81	0.80	0.68
5	Class 5	0.19	0.58	0.17
6	Class 6	0.69	0.67	0.51
7	Class 7	0.79	0.75	0.63
8	Class 8	0.98	0.99	0.97
9	Class 9	0.95	0.97	0.93
10	Class 10	0.93	0.90	0.84
11	Class 11	0.99	0.99	0.99
12	Class 12	0.90	0.84	0.77
13	Class 13	0.63	0.86	0.57

5. Prediction

- ### Predict on Train, Val, Test Dataset

In [41]:

```
def create_mask(pred_mask):  
    pred_mask = tf.argmax(pred_mask, axis=-1)  
    pred_mask = pred_mask[..., tf.newaxis]  
    return pred_mask[0]
```

In [42]:

```
def display(display_list):
    plt.figure(figsize=(15, 15))

    title = ['Input Image', 'True Mask', 'Predicted Mask']

    for i in range(len(display_list)):
        plt.subplot(1, len(display_list), i+1)
        plt.title(title[i])
        plt.imshow(tf.keras.preprocessing.image.array_to_img(display_list[i]))
        plt.axis('off')
    plt.show()
```

In [43]:

```
def show_predictions(dataset, num, model):
    """
    Displays the first image of each of the num batches
    """
    if dataset:
        for image, mask in dataset.take(num):
            pred_mask = model.predict(image)
            display([image[0], mask[0], create_mask(pred_mask)])
    else:
        display([sample_image, sample_mask,
                 create_mask(model.predict(sample_image[tf.newaxis, ...]))])
)
```

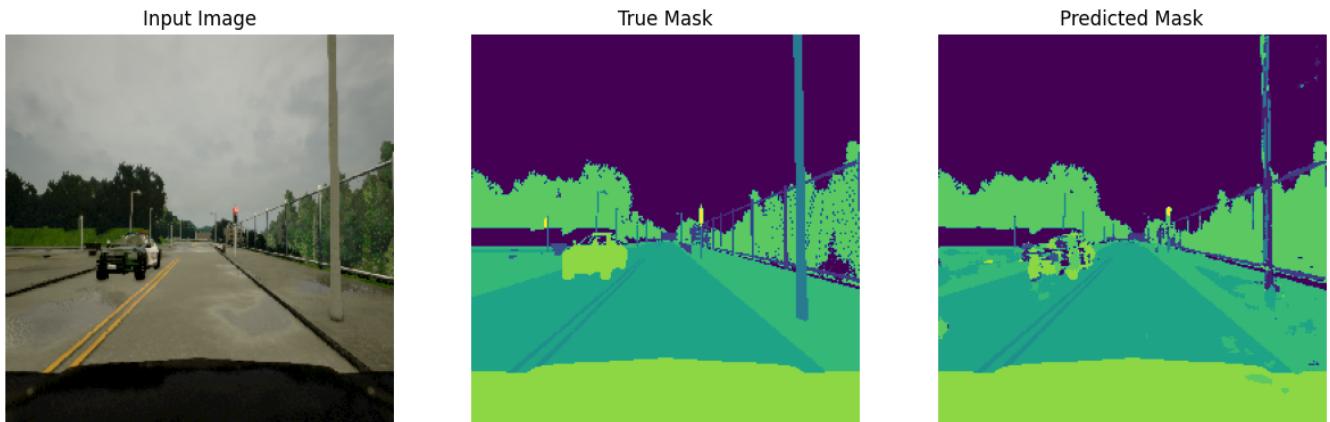
In [44]:

```
# fcn
show_predictions(train_dataset, 3, fcn_model1)
show_predictions(val_dataset, 3, fcn_model1)
show_predictions(test_dataset, 3, fcn_model1)
```

1/1 ————— 0s 53ms/step



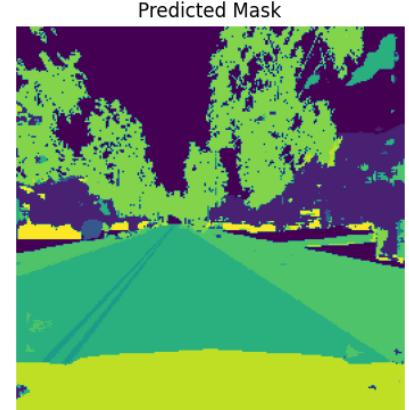
1/1 ————— 0s 36ms/step



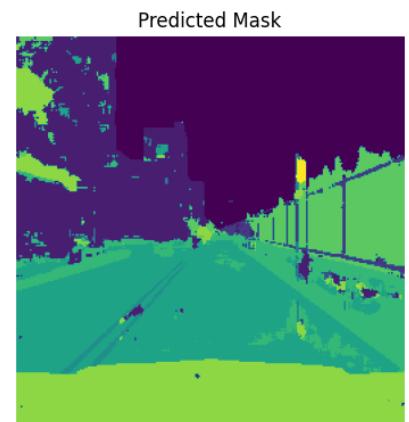
1/1 ————— 0s 29ms/step



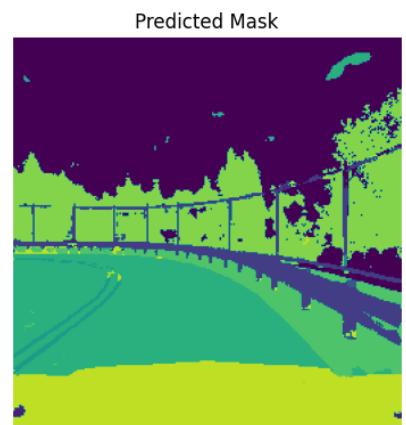
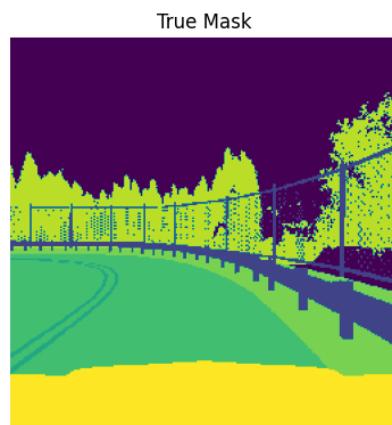
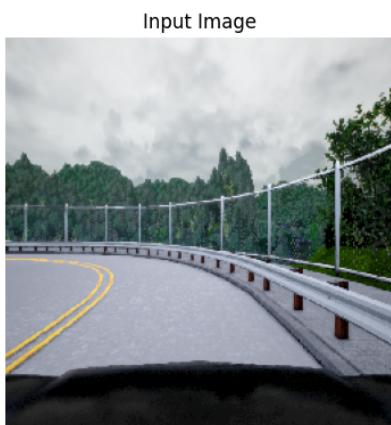
1/1 ————— 0s 46ms/step



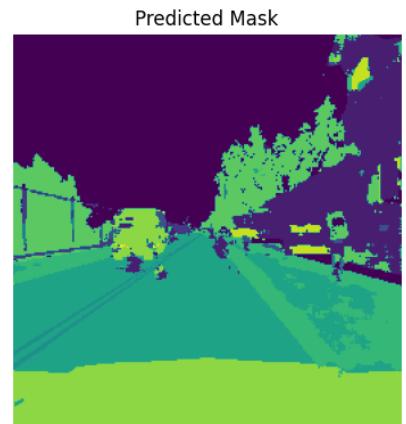
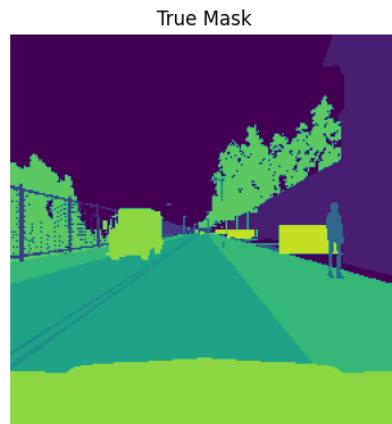
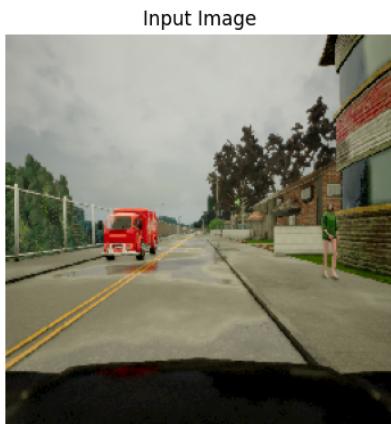
1/1 ————— 0s 53ms/step



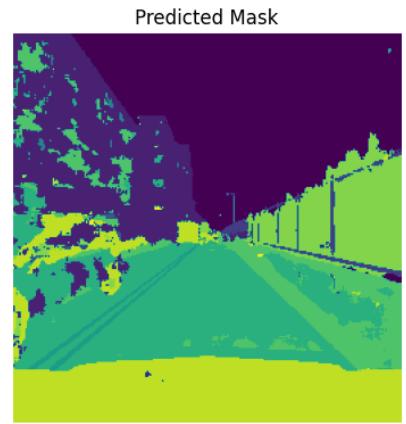
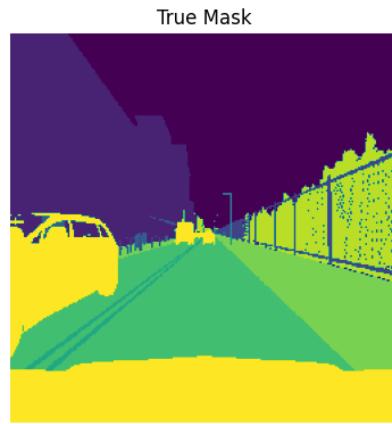
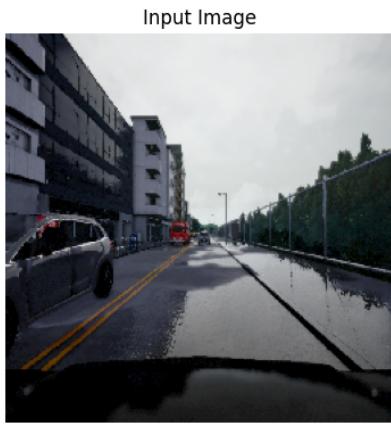
1/1 ————— 0s 28ms/step



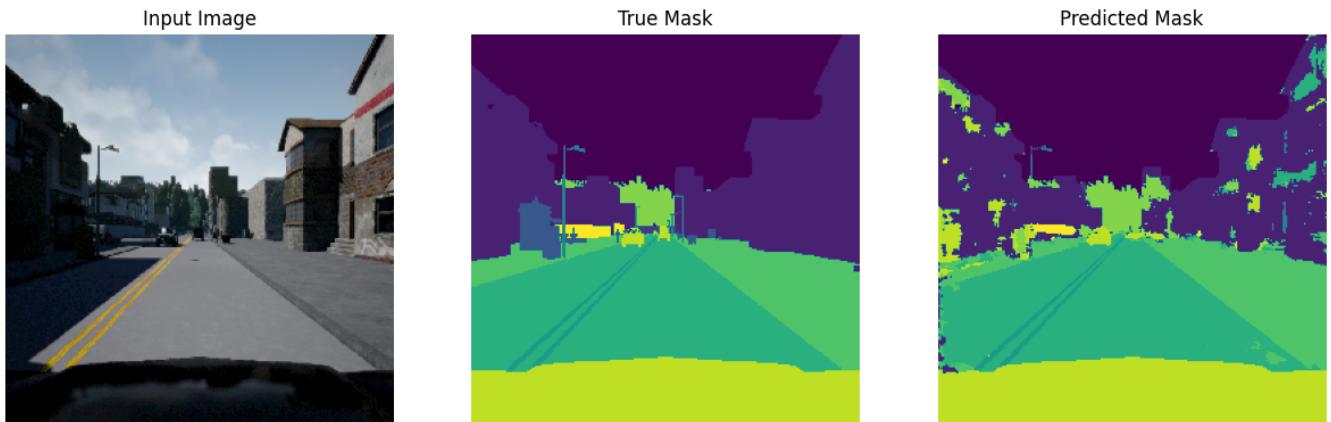
1/1 ————— **0s** 29ms/step



1/1 ————— **0s** 30ms/step



1/1 ————— 0s 29ms/step



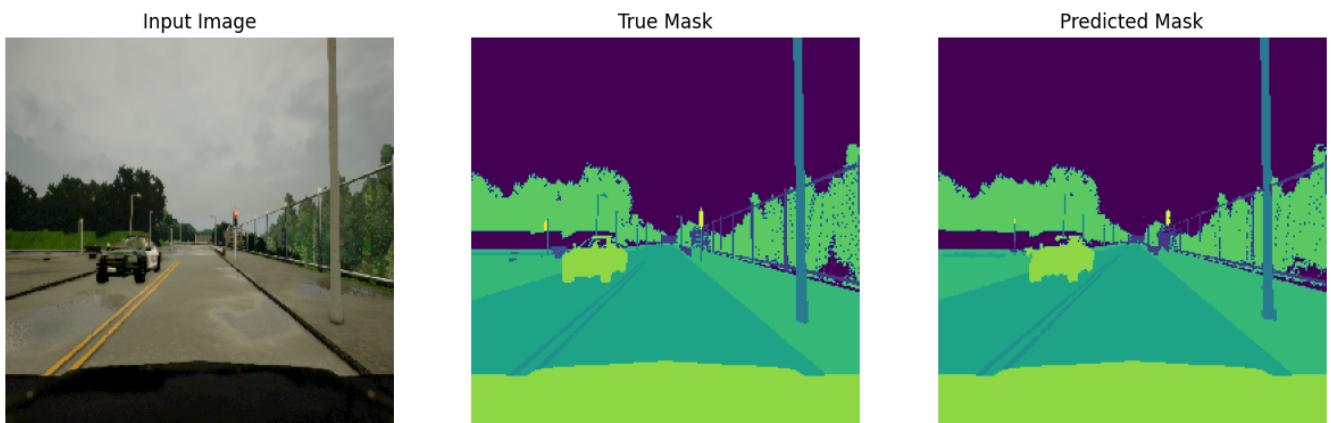
In [45]:

```
# unet  
show_predictions(train_dataset, 3, unet_model1)  
show_predictions(val_dataset, 3, unet_model1)  
show_predictions(test_dataset, 3, unet_model1)
```

1/1 ————— 0s 38ms/step



1/1 ————— 0s 44ms/step



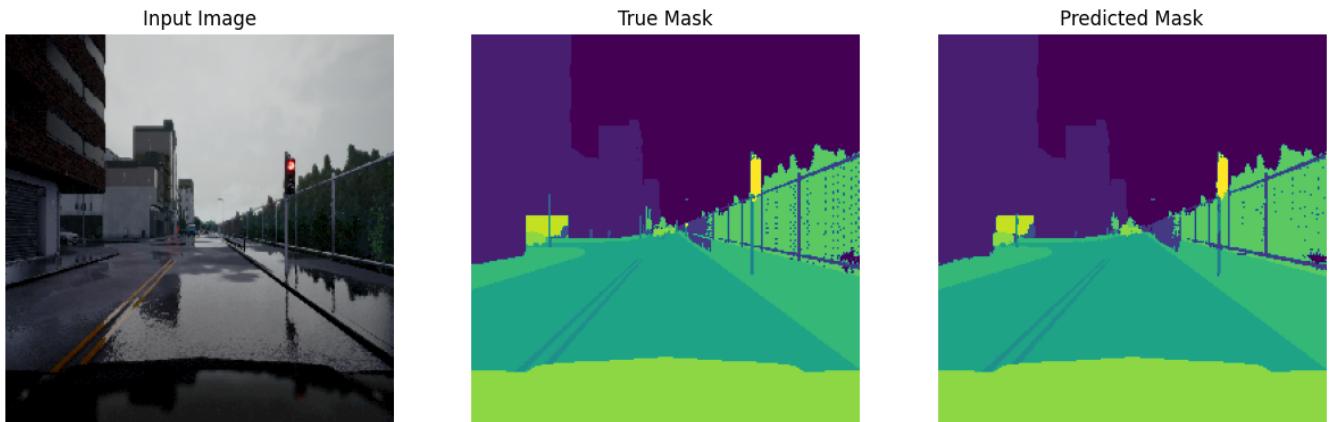
1/1 ————— 0s 30ms/step



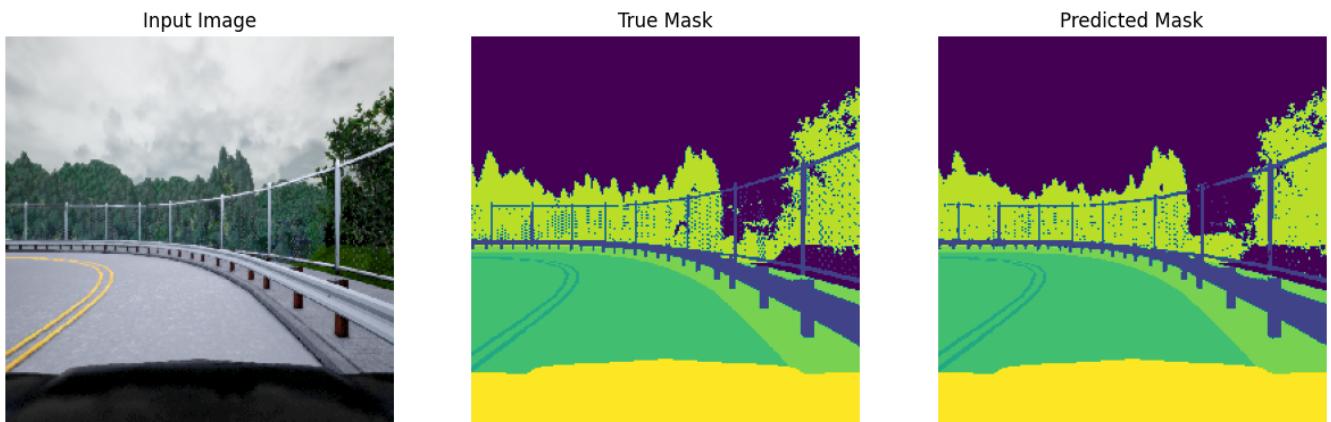
1/1 ————— 0s 38ms/step



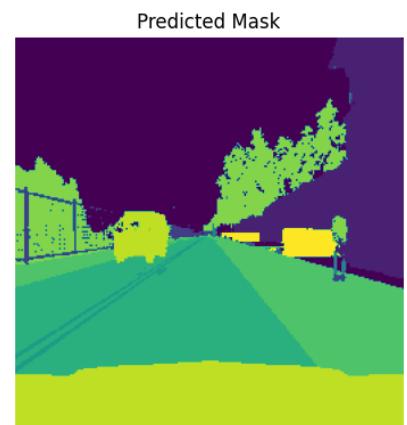
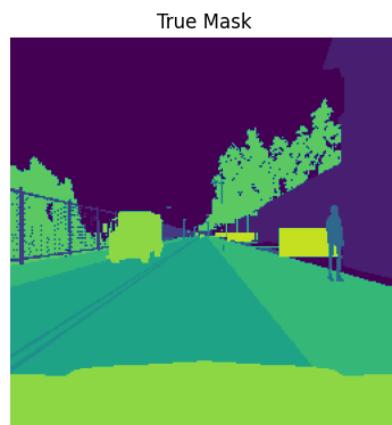
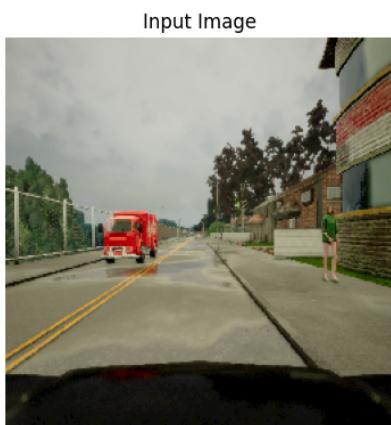
1/1 ————— 0s 38ms/step



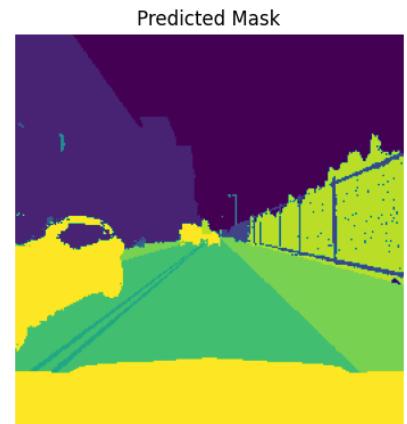
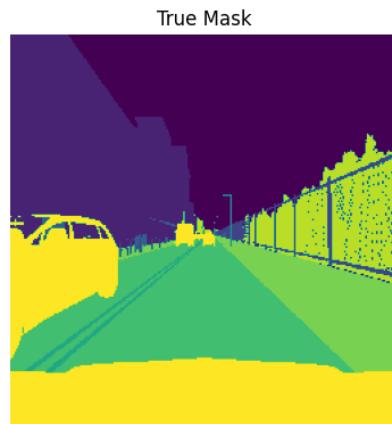
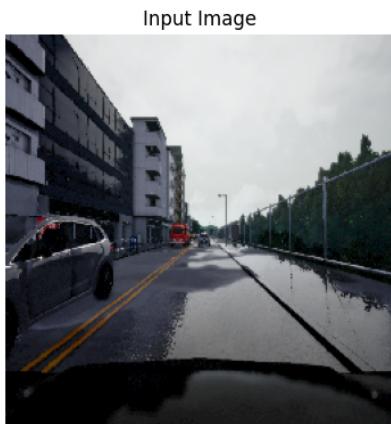
1/1 ————— 0s 30ms/step



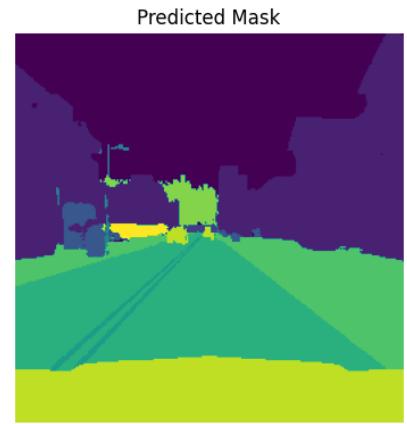
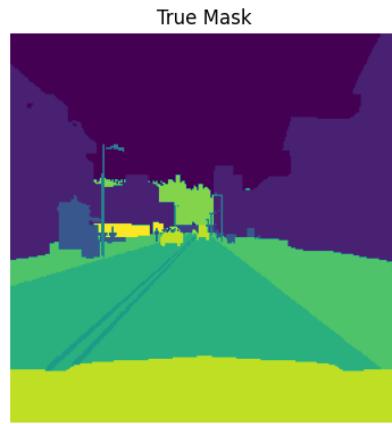
1/1 ————— 0s 30ms/step



1/1 ————— 0s 42ms/step



1/1 ————— 0s 30ms/step



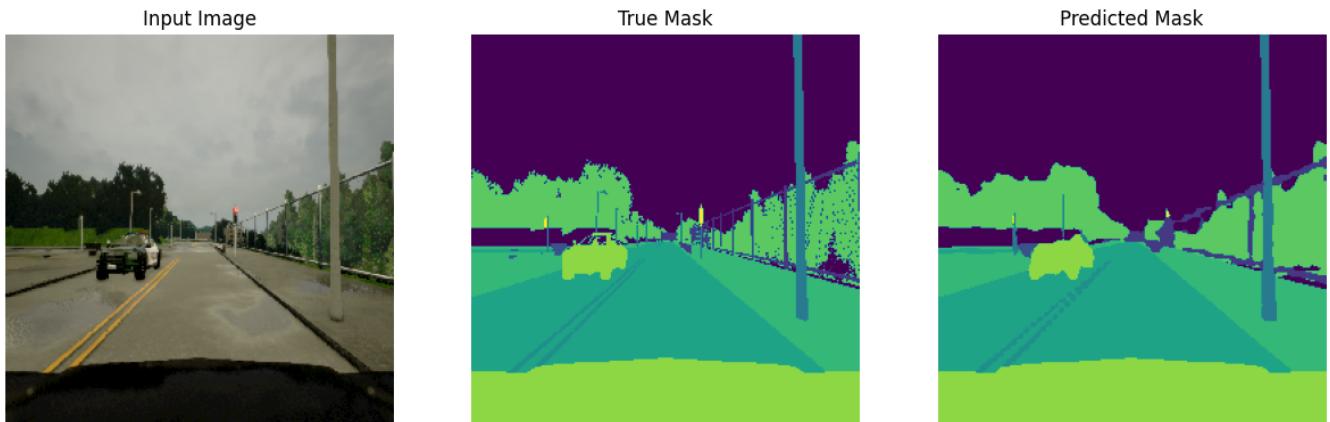
In [46]:

```
# deeplab  
show_predictions(train_dataset, 3, deeplab_model1)  
show_predictions(val_dataset, 3, deeplab_model1)  
show_predictions(test_dataset, 3, deeplab_model1)
```

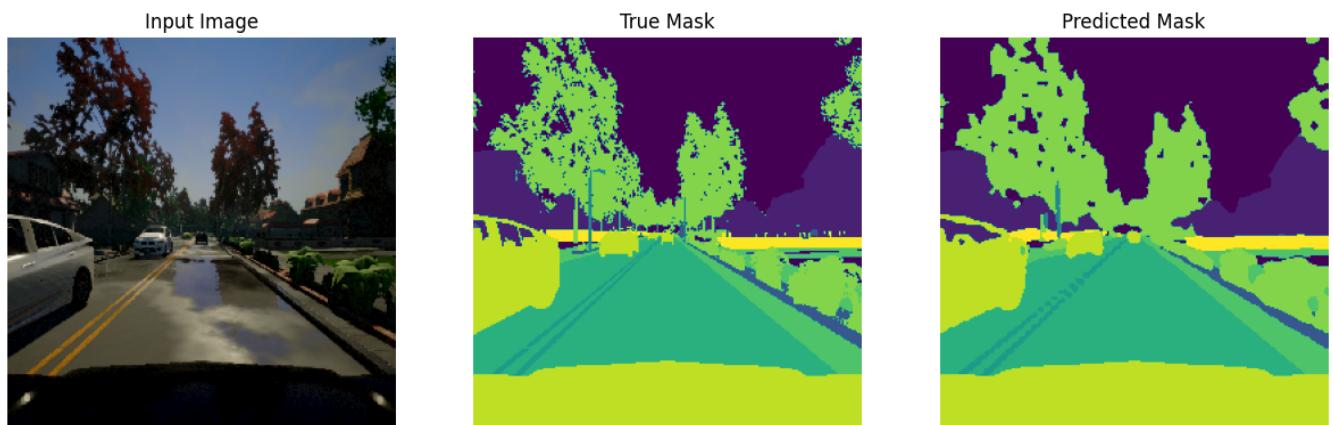
1/1 ————— 0s 54ms/step



1/1 ————— 0s 60ms/step



1/1 ————— 0s 34ms/step



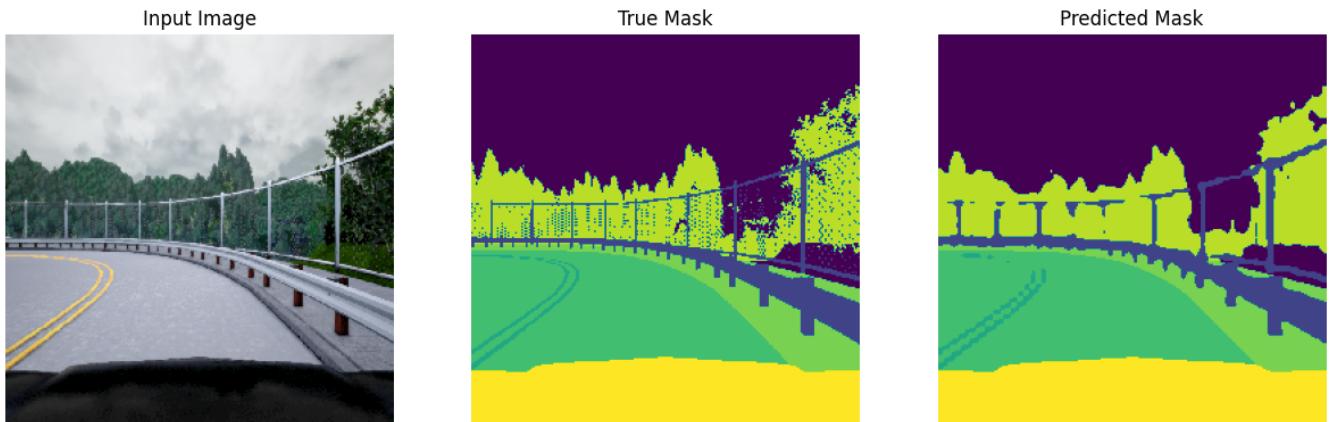
1/1 ————— 0s 76ms/step



1/1 ————— 0s 72ms/step



1/1 ————— 0s 34ms/step



1/1 ————— 0s 46ms/step



1/1 ————— 0s 35ms/step



1/1 ————— **0s** 32ms/step



- ### Predict on New Dataset

In [47]:

```
from pathlib import Path

def load_and_preprocess_image(image_path, mask_path):
    image = tf.io.read_file(image_path)
    image = tf.image.decode_png(image, channels=3)
    image = tf.image.resize(image, [256, 256], method='nearest')
    image = tf.image.convert_image_dtype(image, tf.float32)

    mask = tf.io.read_file(mask_path)
    mask = tf.image.decode_png(mask, channels=3)
    mask = tf.image.resize(mask, [256, 256], method='nearest')
    mask = tf.math.reduce_max(mask, axis=-1, keepdims=True)

    return image, mask

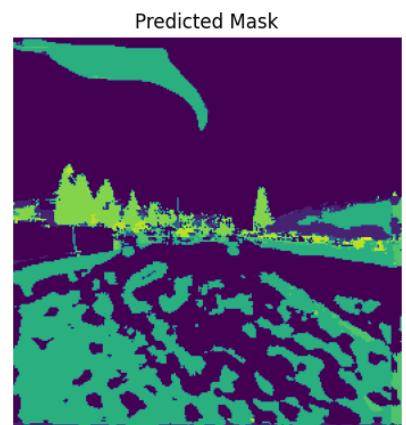
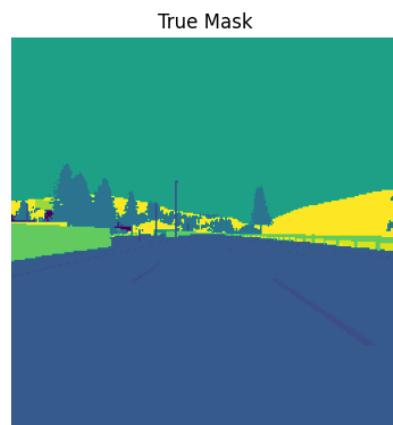
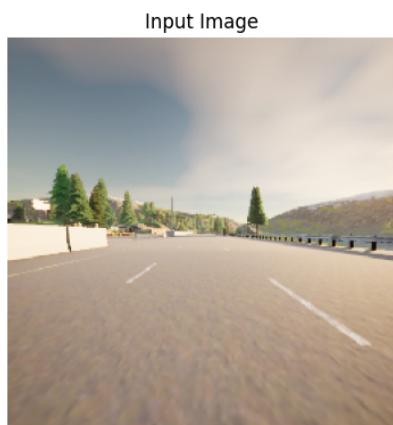
# Paths to the new dataset
IMAGE_PATH = Path("/kaggle/input/semantic-segmentation-car-driving/x-rgb")
MASK_PATH = Path("/kaggle/input/semantic-segmentation-car-driving/y-mask")
new_images_paths = sorted([str(x) for x in IMAGE_PATH.glob("*.png")])
new_masks_paths = sorted([str(x) for x in MASK_PATH.glob("*.png")])

# Creating the new dataset
new_dataset = tf.data.Dataset.from_tensor_slices((new_images_paths, new_masks_paths))
new_dataset = new_dataset.map(load_and_preprocess_image)
new_dataset_batched = new_dataset.batch(1)
```

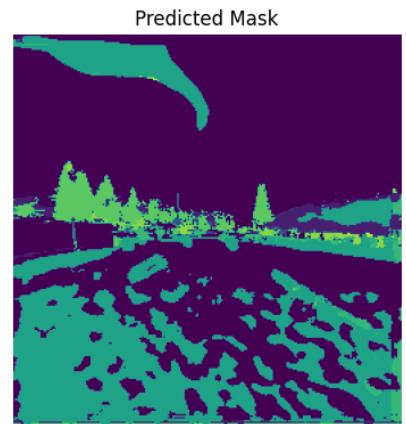
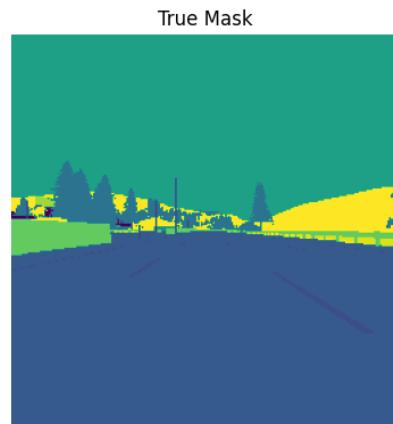
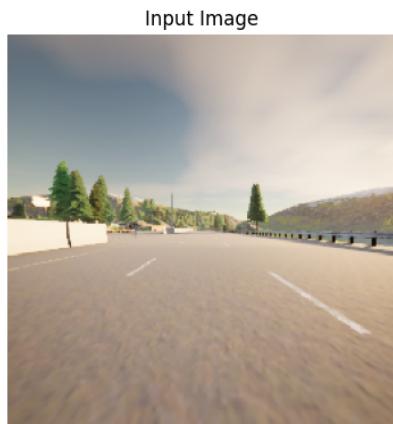
In [48]:

```
# fcn
show_predictions(new_dataset_batched, 5, fcn_model1)
```

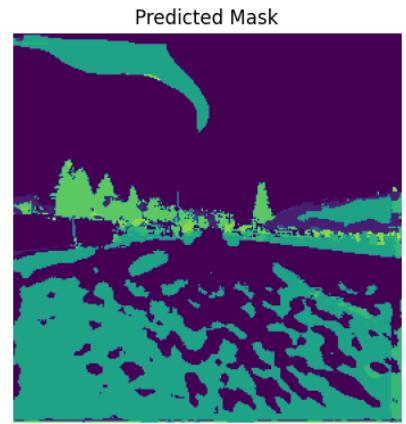
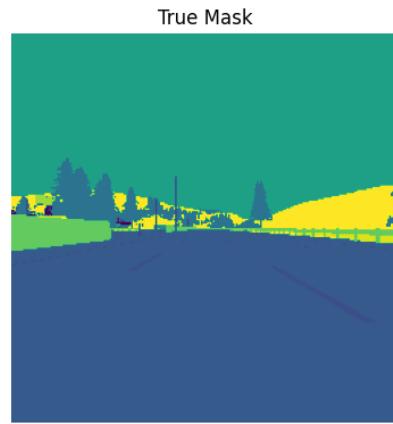
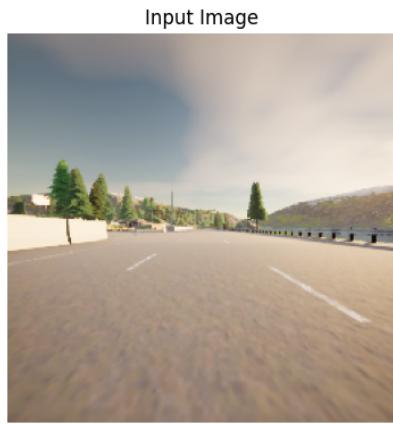
1/1 ————— 1s 1s/step



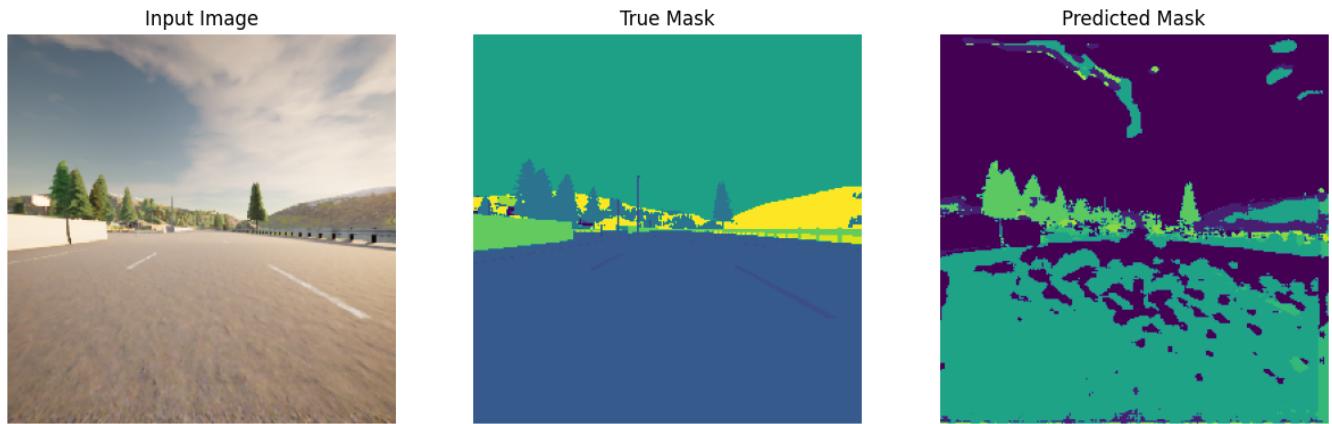
1/1 ————— 0s 18ms/step



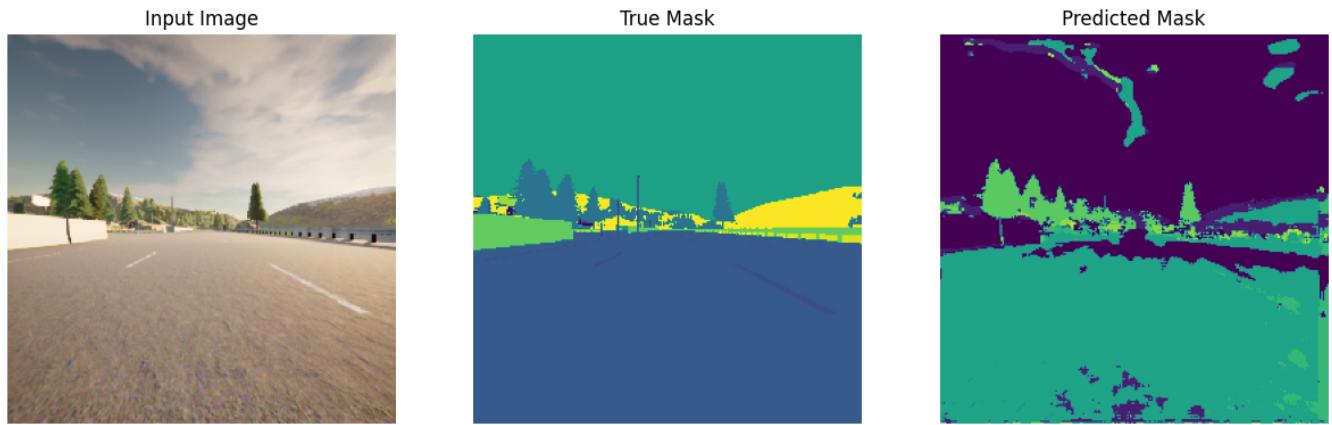
1/1 ————— 0s 18ms/step



1/1 ————— **0s** 18ms/step



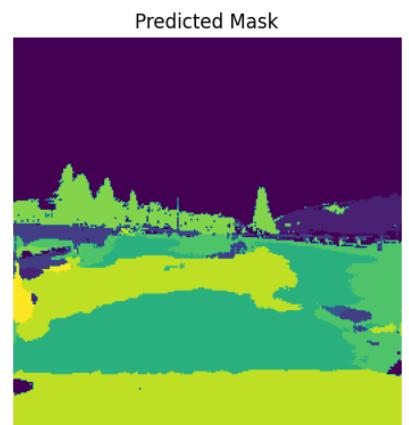
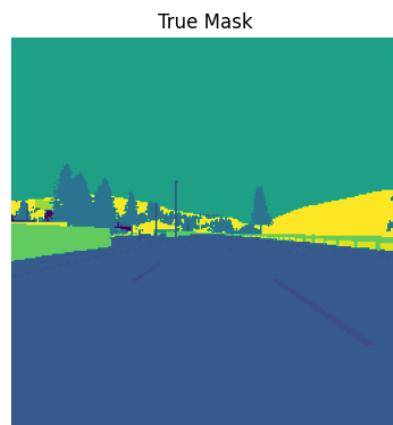
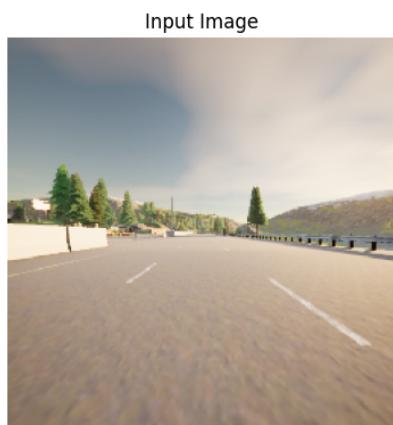
1/1 ————— **0s** 18ms/step



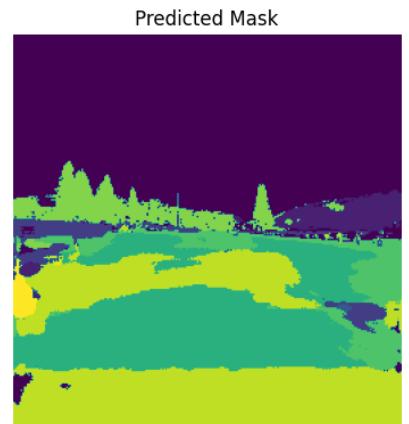
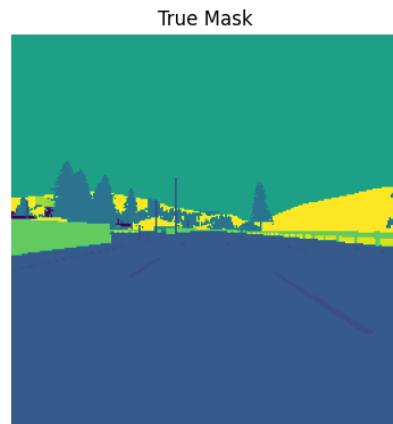
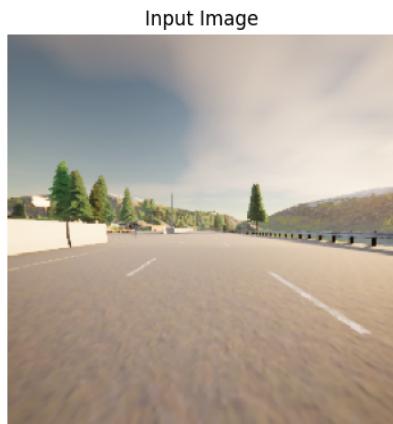
In [49]:

```
# unet
show_predictions(new_dataset_batched, 5, unet_model1)
```

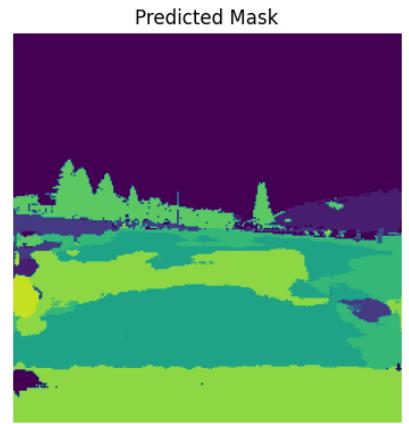
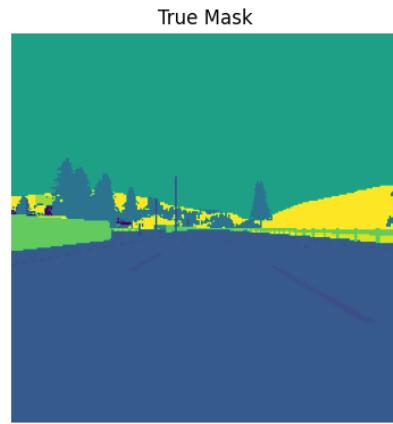
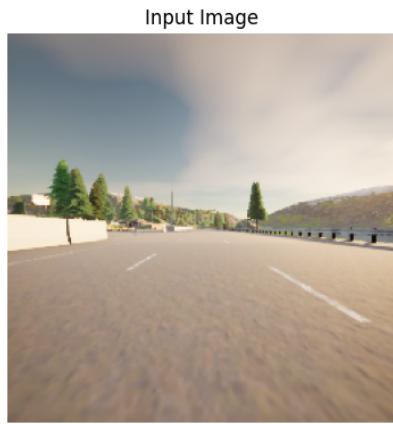
1/1 ————— **2s** 2s/step



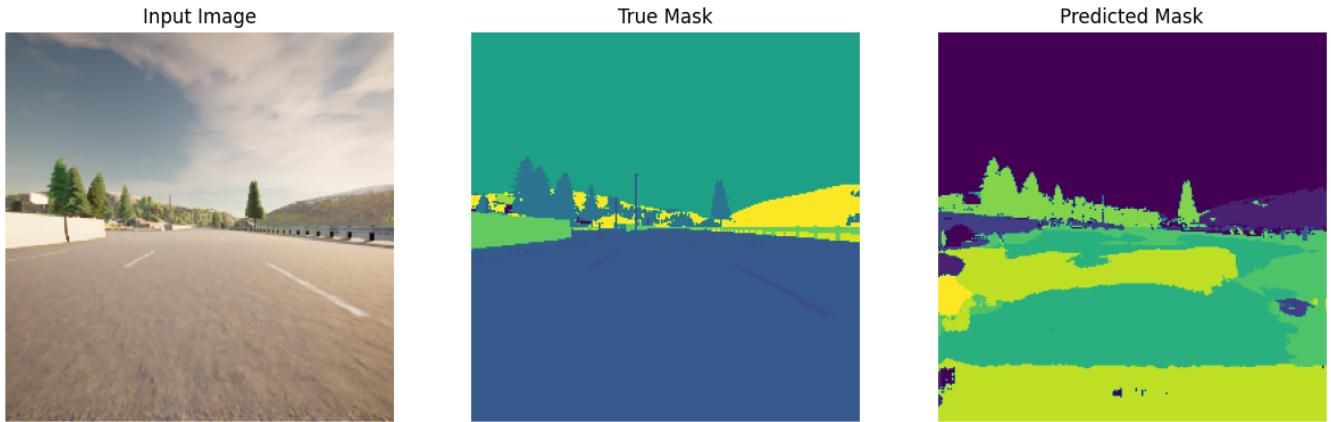
1/1 ————— **0s** 19ms/step



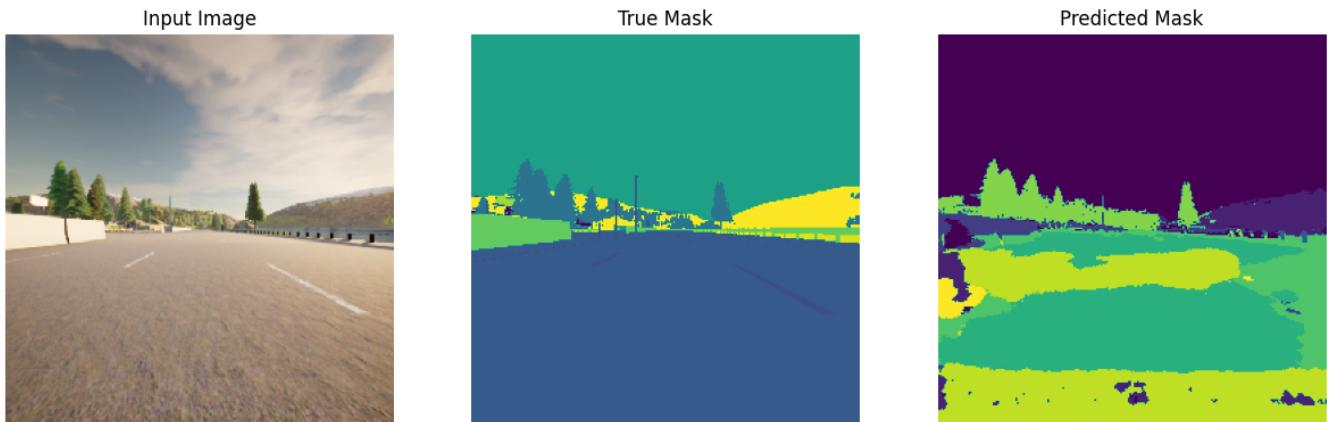
1/1 ————— **0s** 19ms/step



1/1 ————— **0s** 19ms/step



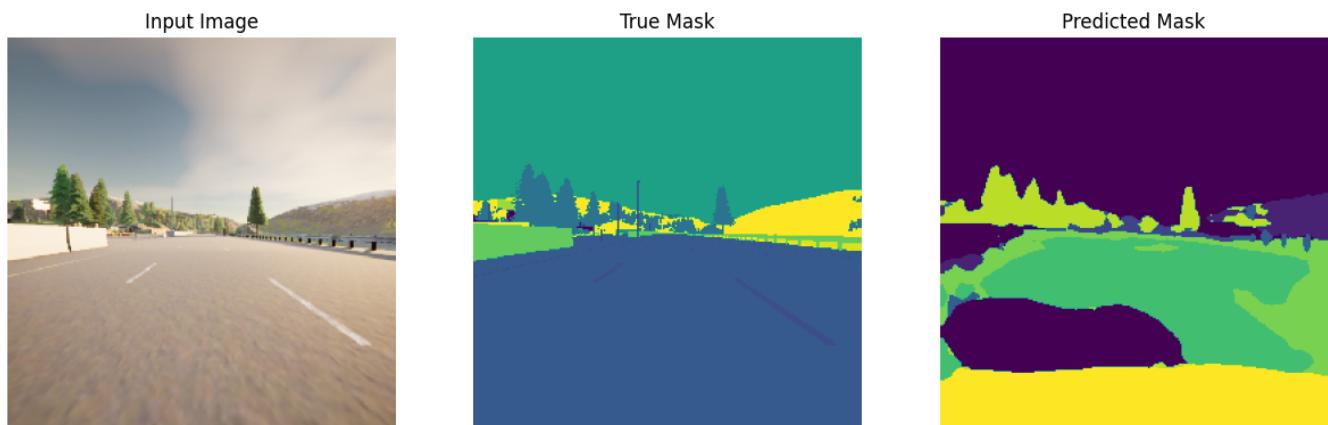
1/1 ————— **0s** 19ms/step



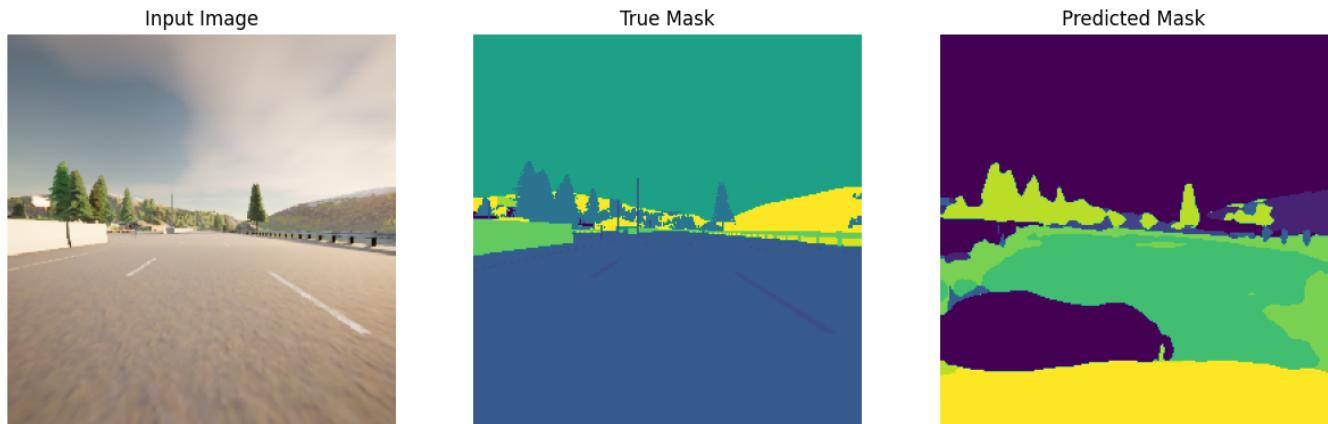
In [50]:

```
#deeplab  
show_predictions(new_dataset_batched, 5, deeplab_model1)
```

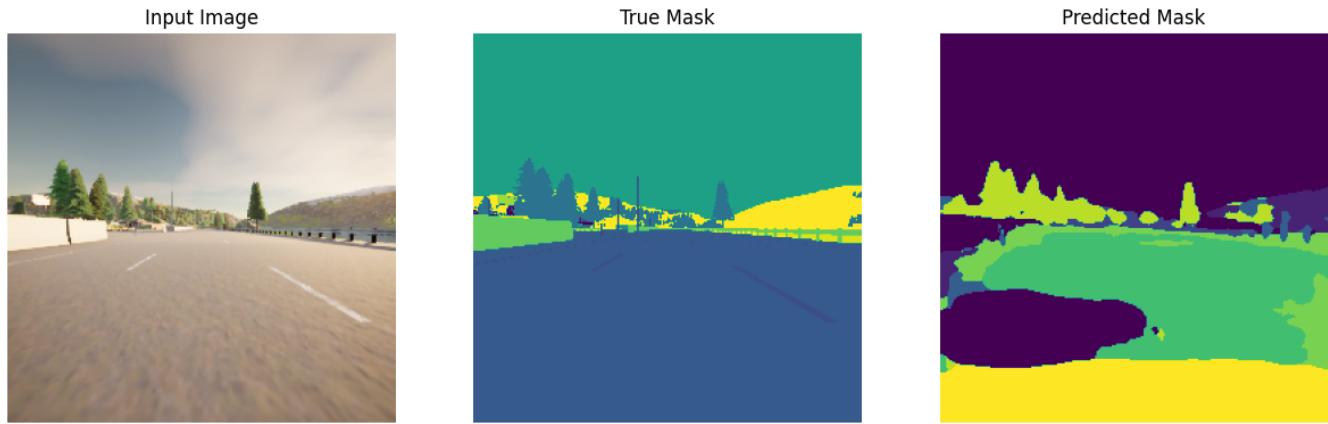
1/1 ————— **4s** 4s/step



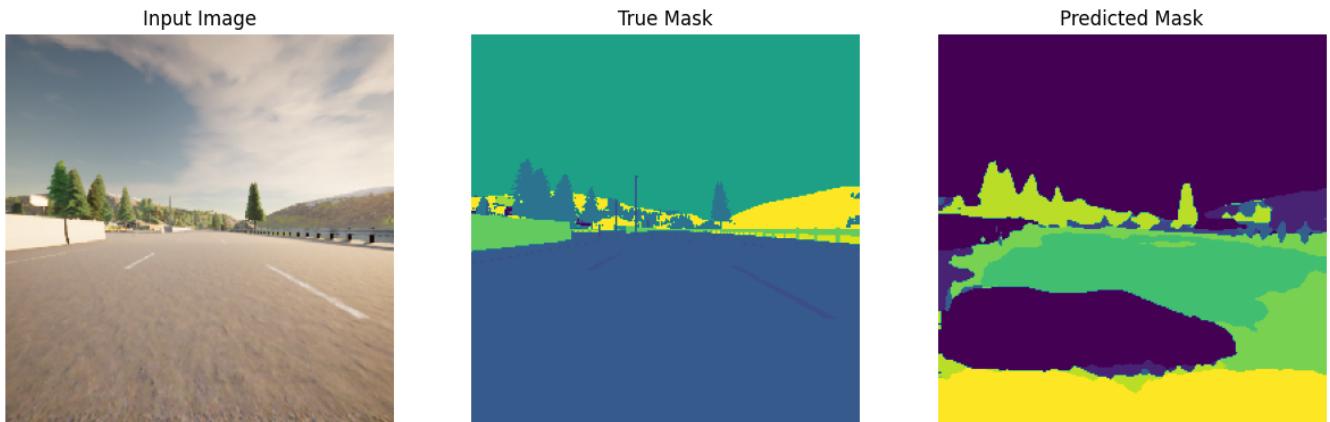
1/1 ————— 0s 22ms/step



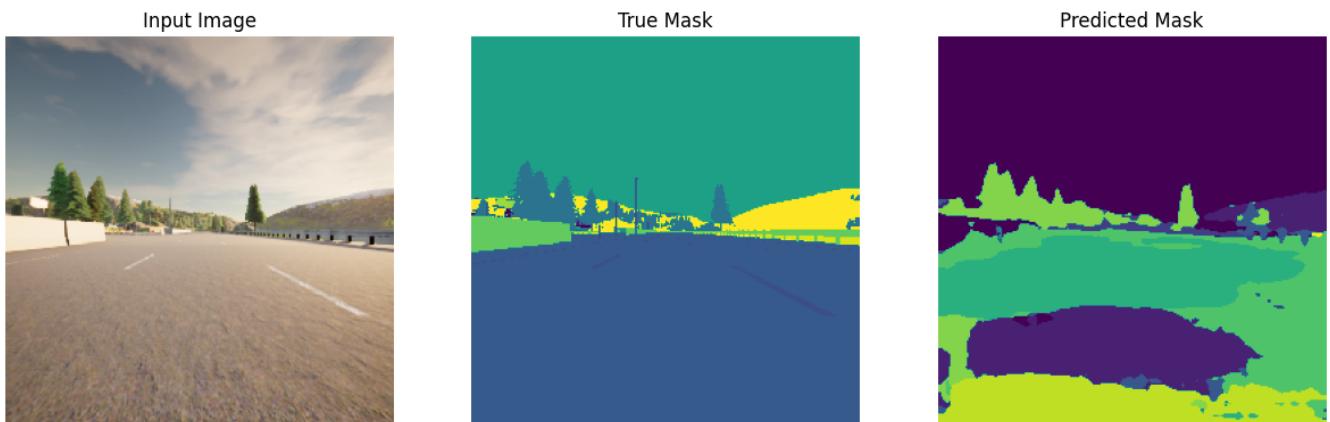
1/1 ————— 0s 22ms/step



1/1 ————— 0s 22ms/step



1/1 ————— 0s 22ms/step



6. Save Model

In [51]:

```
from tensorflow.keras.models import load_model
fcn_model1.save('fcn1-image-segmentation-model.h5')

unet_model1.save('unet1-image-segmentation-model.h5')

deeplab_model1.save('deeplab1-image-segmentation-model.h5')
```